MAX POOLING

Input

$$\begin{bmatrix} 6 & 3 & 2 \\ 5 & 2 & 1 \\ 4 & 1 & 0 \end{bmatrix}$$

Pooling matrix 2x2

Output

$$\begin{bmatrix} 6 & 2 \\ 4 & 0 \end{bmatrix}$$

Pytorch.ipynb - Colab | CNN Explainer | Animations of Convolution and | Image Kernels explained visuall

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=sud58LClW23n

Paused

+ Code + Text All changes saved

RAM Disk

Convolution using Pytorch

2s

```
import torch
import torch.nn as nn
from PIL import Image
import numpy as np
import matplotlib.pyplot as plt

# Load and preprocess the image
def load_image(image_path, image_size=(128, 128)):
    image = Image.open(image_path).convert('L')  # Convert to grayscale
    # image = image.resize(image_size)             # Resize the image
    image = np.array(image)                      # Convert to numpy array
    image = torch.tensor(image, dtype=torch.float32)  # Convert to torch tensor
    image = image.unsqueeze(0).unsqueeze(0)      # Add batch and channel dimensions
    return image

# Path to the image
image_path = r"/content/drive/MyDrive/Dataset/images.jpeg"

# Load the image
input_image = load_image(image_path)
print("Input Image Shape:", input_image.shape)

# Define a convolutional layer
conv_layer = nn.Conv2d(in_channels=1, out_channels=1, kernel_size=3, stride=1, padding=0)
```

1s completed at 9:22 PM

Pytorch.ipynb - Colab | CNN Explainer | Animations of Convolution and | Image Kernels explained visuall

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=sud58LClW23n

Paused

+ Code + Text All changes saved

RAM Disk

```
array([[0.09865053, 0.52521243, 0.36051894, 0.21917541, 0.96039324],
       [0.36312466, 0.20557749, 0.09425741, 0.21215107, 0.04927506],
       [0.7737655 , 0.10031832, 0.44264185, 0.97337431, 0.33021284],
       [0.11717513, 0.63664965, 0.31029961, 0.4635026 , 0.76889252],
       [0.45942061, 0.89752827, 0.7664678 , 0.72507998, 0.83271031]])
```

Convolution using Pytorch

```python
import torch
import torch.nn as nn
from PIL import Image
import numpy as np
import matplotlib.pyplot as plt

# Load and preprocess the image
def load_image(image_path, image_size=(128, 128)):
    image = Image.open(image_path).convert('L')  # Convert to grayscale
    # image = image.resize(image_size)             # Resize the image
    image = np.array(image)                      # Convert to numpy array
    image = torch.tensor(image, dtype=torch.float32)  # Convert to torch tensor
    image = image.unsqueeze(0).unsqueeze(0)      # Add batch and channel dimensions
    return image

# Path to the image
image_path = r"/content/drive/MyDrive/Dataset/images.jpeg"

# Load the image
```

1s completed at 9:22 PM

Pytorch.ipynb - Colab | CNN Explainer | Animations of Convolution and | Image Kernels explained visuall

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=Kbscvi23wKvu

Paused

+ Code + Text

RAM Disk

[2] 0s

```
array([[0.09865053, 0.52521243, 0.36051894, 0.21917541, 0.96039324],
       [0.36312466, 0.20557749, 0.09425741, 0.21215107, 0.04927506],
       [0.7737655 , 0.10031832, 0.44264185, 0.97337431, 0.33021284],
       [0.11717513, 0.63664965, 0.31029961, 0.4635026 , 0.76889252],
       [0.45942061, 0.89752827, 0.7664678 , 0.72507998, 0.83271031]])
```

Convolution and Pooling using Pytorch

Convolution and using Pytorch

[8] 2s

```python
import torch
import torch.nn as nn
from PIL import Image
import numpy as np
import matplotlib.pyplot as plt

# Load and preprocess the image
def load_image(image_path, image_size=(128, 128)):
    image = Image.open(image_path).convert('L')  # Convert to grayscale
    # image = image.resize(image_size)             # Resize the image
    image = np.array(image)                      # Convert to numpy array
    image = torch.tensor(image, dtype=torch.float32)  # Convert to torch tensor
    image = image.unsqueeze(0).unsqueeze(0)      # Add batch and channel dimensions
    return image
```

✓ 1s completed at 9:22 PM

AISCIENCES

```
    image = image.unsqueeze(0).unsqueeze(0)      # Add batch and channel dimensions
    return image

# Path to the image
image_path = r"/content/drive/MyDrive/Dataset/images.jpeg"

# Load the image
input_image = load_image(image_path)
print("Input Image Shape:", input_image.shape)

# Define a convolutional layer
conv_layer = nn.Conv2d(in_channels=1, out_channels=1, kernel_size=3, stride=1, padding=0)

# Set a specific filter for the convolutional layer
conv_layer.weight = nn.Parameter(torch.tensor([[[[1, 0, -1],
                                                 [1, 0, -1],
                                                 [1, 0, -1]]]], dtype=torch.float32))

# Perform convolution
conv_output = conv_layer(input_image)
max_pooling


# Display the image
plt.imshow(input_image[0, 0].numpy(), cmap='gray')
plt.title("Input Image")
plt.show()
print("Input Image Shape:", input_image.shape)
```

Pytorch.ipynb - Colab | CNN Explainer | Animations of Convolution and | Image Kernels explained visuall

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=sud58LClW23n

Paused

+ Code + Text All changes saved

RAM Disk

```
    image = image.unsqueeze(0).unsqueeze(0)      # Add batch and channel dimensions
    return image

# Path to the image
image_path = r"/content/drive/MyDrive/Dataset/images.jpeg"

# Load the image
input_image = load_image(image_path)
print("Input Image Shape:", input_image.shape)

# Define a convolutional layer
conv_layer = nn.Conv2d(in_channels=1, out_channels=1, kernel_size=3, stride=1, padding=0)

# Set a specific filter for the convolutional layer
conv_layer.weight = nn.Parameter(torch.tensor([[[[1, 0, -1],
                                                 [1, 0, -1],
                                                 [1, 0, -1]]]], dtype=torch.float32))

# Perform convolution
conv_output = conv_layer(input_image)
max_pooling_img = nn.MaxPool2d(kernel_size=2, stride=2)


# Display the image
plt.imshow(input_image[0, 0].numpy(), cmap='gray')
plt.title("Input Image")
plt.show()
print("Input Image Shape:", input_image.shape)
```

1s completed at 9:22 PM

Pytorch.ipynb - Colab | CNN Explainer | Animations of Convolution and | Image Kernels explained visuall

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=sud58LClW23n

Paused

+ Code + Text All changes saved

RAM Disk

```
# Define a convolutional layer
conv_layer = nn.Conv2d(in_channels=1, out_channels=1, kernel_size=3, stride=1, padding=0)

# Set a specific filter for the convolutional layer
conv_layer.weight = nn.Parameter(torch.tensor([[[[1, 0, -1],
                                                 [1, 0, -1],
                                                 [1, 0, -1]]]], dtype=torch.float32))

# Perform convolution
conv_output = conv_layer(input_image)
max_pooling_img = nn.MaxPool2d(kernel_size=2, stride=2)


# Display the image
plt.imshow(input_image[0, 0].numpy(), cmap='gray')
plt.title("Input Image")
plt.show()
print("Input Image Shape:", input_image.shape)

# Display the image
plt.imshow(conv_output[0, 0].detach().numpy(), cmap='gray')
plt.title("Conv Image")
plt.show()
print("Conv Image Shape:", conv_output.shape)
```

```
Input Image Shape: torch.Size([1, 1, 225, 225])
```

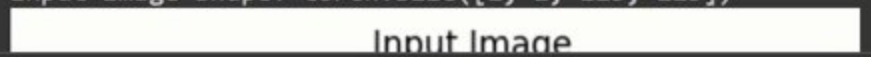

1s completed at 9:22 PM

AISCIENCES

Pytorch.ipynb - Colab | CNN Explainer | Animations of Convolution and | Image Kernels explained visuall

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=sud58LClW23n

Paused

+ Code + Text Saving...

RAM Disk

```
# Set a specific filter for the convolutional layer
conv_layer.weight = nn.Parameter(torch.tensor([[[[1, 0, -1],
                                                 [1, 0, -1],
                                                 [1, 0, -1]]]], dtype=torch.float32))

# Perform convolution
conv_output = conv_layer(input_image)
max_pooling_img = nn.MaxPool2d(kernel_size=2, stride=2)


# Display the image
plt.imshow(input_image[0, 0].numpy(), cmap='gray')
plt.title("Input Image")
plt.show()
print("Input Image Shape:", input_image.shape)

# Display the image
plt.imshow(conv_output[0, 0].detach().numpy(), cmap='gray')
plt.title("Conv Image")
plt.show()
print("Conv Image Shape:", conv_output.shape)

# Display the image
plt.imshow(conv_output[0, 0].detach().numpy(), cmap='gray')
plt.title("Conv Image")
plt.show()
print("Conv Image Shape:", conv_output.shape)
```

1s completed at 9:22 PM

AISCIENCES

Pytorch.ipynb - Colab | CNN Explainer | Animations of Convolution and | Image Kernels explained visuall

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=sud58LClW23n

Paused

+ Code + Text

RAM Disk

```
# Set a specific filter for the convolutional layer
conv_layer.weight = nn.Parameter(torch.tensor([[[[1, 0, -1],
                                                 [1, 0, -1],
                                                 [1, 0, -1]]]], dtype=torch.float32))

# Perform convolution
conv_output = conv_layer(input_image)
max_pooling_img = nn.MaxPool2d(kernel_size=2, stride=2)


# Display the image
plt.imshow(input_image[0, 0].numpy(), cmap='gray')
plt.title("Input Image")
plt.show()
print("Input Image Shape:", input_image.shape)

# Display the image
plt.imshow(conv_output[0, 0].detach().numpy(), cmap='gray')
plt.title("Conv Image")
plt.show()
print("Conv Image Shape:", conv_output.shape)

# Display the image
plt.imshow(Pooling[0, 0].detach().numpy(), cmap='gray')
plt.title("Pooling Image")
plt.show()
print("Pooling Image Shape:", conv_output.shape)
```

1s completed at 9:22 PM

Pytorch.ipynb - Colab | CNN Explainer | Animations of Convolution and | Image Kernels explained visuall

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=sud58LClW23n

Paused

+ Code + Text

RAM Disk

```
# Set a specific filter for the convolutional layer
conv_layer.weight = nn.Parameter(torch.tensor([[[[1, 0, -1],
                                                 [1, 0, -1],
                                                 [1, 0, -1]]]], dtype=torch.float32))

# Perform convolution
conv_output = conv_layer(input_image)
max_pooling_img = nn.MaxPool2d(kernel_size=2, stride=2)


# Display the image
plt.imshow(input_image[0, 0].numpy(), cmap='gray')
plt.title("Input Image")
plt.show()
print("Input Image Shape:", input_image.shape)

# Display the image
plt.imshow(conv_output[0, 0].detach().numpy(), cmap='gray')
plt.title("Conv Image")
plt.show()
print("Conv Image Shape:", conv_output.shape)

# Display the image
plt.imshow(max_pooling_img[0, 0].detach().numpy(), cmap='gray')
plt.title("Pooling Image")
plt.show()
print("Pooling Image Shape:", max_pooling_img.shape)
```

✓ 1s completed at 9:22 PM

Pytorch.ipynb - Colab | CNN Explainer | Animations of Convolution and | Image Kernels explained visuall

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=sud58LClW23n

Paused

+ Code + Text Saving...

RAM Disk

```
    return image

# Path to the image
image_path = r"/content/drive/MyDrive/Dataset/images.jpeg"

# Load the image
input_image = load_image(image_path)
print("Input Image Shape:", input_image.shape)

# Define a convolutional layer
conv_layer = nn.Conv2d(in_channels=1, out_channels=1, kernel_size=3, stride=1, padding=0)

# Set a specific filter for the convolutional layer
conv_layer.weight = nn.Parameter
```

```
(kernel_size: _size_any_t, stride: _size_any_t | None =
None, padding: _size_any_t = 0, dilation: _size_any_t =
1, return_indices: bool = False, ceil_mode: bool =
False) -> None

Initialize internal Module state, shared by both nn.Module and ScriptM
```

```
# Perform convolution
conv_output = conv_layer(input_i
max_pooling_img = nn.MaxPool2d(c, kernel_size=2, stride=2)


# Display the image
plt.imshow(input_image[0, 0].numpy(), cmap='gray')
plt.title("Input Image")
plt.show()
print("Input Image Shape:", input_image.shape)
```

0s completed at 9:41 PM

AISCIENCES

Pytorch.ipynb - Colab | CNN Explainer | Animations of Convolution and | Image Kernels explained visuall

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=sud58LClW23n

Paused

+ Code + Text All changes saved

RAM Disk

```
    return image

# Path to the image
image_path = r"/content/drive/MyDrive/Dataset/images.jpeg"

# Load the image
input_image = load_image(image_path)
print("Input Image Shape:", input_image.shape)

# Define a convolutional layer
conv_layer = nn.Conv2d(in_channels=1, out_channels=1, kernel_size=3, stride=1, padding=0)

# Set a specific filter for the convolutional layer
conv_layer.weight = nn.Parameter(torch.tensor([[[[1, 0, -1],
                                                 [1, 0, -1],
                                                 [1, 0, -1]]]], dtype=torch.float32))

# Perform convolution
conv_output = conv_layer(input_image)
max_pooling_img = nn.MaxPool2d(conv_output, kernel_size=2, stride=2)


# Display the image
plt.imshow(input_image[0, 0].numpy(), cmap='gray')
plt.title("Input Image")
plt.show()
print("Input Image Shape:", input_image.shape)
```

0s completed at 9:42 PM

AISCIENCES

Pytorch.ipynb - Colab | CNN Explainer | Animations of Convolution and | Image Kernels explained visuall

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=sud58LClW23n

Paused

+ Code + Text Saving...

RAM Disk

```
    return image

# Path to the image
image_path = r"/content/drive/MyDrive/Dataset/images.jpeg"

# Load the image
input_image = load_image(image_path)
print("Input Image Shape:", input_image.shape)

# Define a convolutional layer
conv_layer = nn.Conv2d(in_channels=1, out_channels=1, kernel_size=3, stride=1, padding=0)

# Set a specific filter for the convolutional layer
conv_layer.weight = nn.Parameter(torch.tensor([[[[1, 0, -1],
                                                 [1, 0, -1],
                                                 [1, 0, -1]]]], dtype=torch.float32))

# Perform convolution
conv_output = conv_layer(input_image)
max_pooling_layer = nn.MaxPool2d(conv_output, kernel_size=2, stride=2)
max_pooling_img

# Display the image
plt.imshow(input_image[0, 0].numpy(), cmap='gray')
plt.title("Input Image")
plt.show()
print("Input Image Shape:", input_image.shape)
```

0s completed at 9:42 PM

AISCIENCES

Pytorch.ipynb - Colab | CNN Explainer | Animations of Convolution and | Image Kernels explained visuall

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=sud58LClW23n

Paused

+ Code + Text All changes saved

RAM Disk

```
    return image

# Path to the image
image_path = r"/content/drive/MyDrive/Dataset/images.jpeg"

# Load the image
input_image = load_image(image_path)
print("Input Image Shape:", input_image.shape)

# Define a convolutional layer
conv_layer = nn.Conv2d(in_channels=1, out_channels=1, kernel_size=3, stride=1, padding=0)

# Set a specific filter for the convolutional layer
conv_layer.weight = nn.Parameter(torch.tensor([[[[1, 0, -1],
                                                 [1, 0, -1],
                                                 [1, 0, -1]]]], dtype=torch.float32))

# Perform convolution
conv_output = conv_layer(input_image)
max_pooling_layer = nn.MaxPool2d(kernel_size=2, stride=2)
max_pooling_img = max_pooling_layer(conv_output)

# Display the image
plt.imshow(input_image[0, 0].numpy(), cmap='gray')
plt.title("Input Image")
plt.show()
print("Input Image Shape:", input_image.shape)
```

0s completed at 9:42 PM

Pytorch.ipynb - Colab | CNN Explainer | Animations of Convolution and | Image Kernels explained visuall

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=sud58LClW23n

Paused

+ Code + Text All changes saved

RAM Disk

```
plt.imshow(input_image[0, 0].numpy(), cmap='gray')
plt.title("Input Image")
plt.show()
print("Input Image Shape:", input_image.shape)

# Display the image
plt.imshow(conv_output[0, 0].detach().numpy(), cmap='gray')
plt.title("Conv Image")
plt.show()
print("Conv Image Shape:", conv_output.shape)

# Display the image
plt.imshow(max_pooling_img[0, 0].detach().numpy(), cmap='gray')
plt.title("Pooling Image")
plt.show()
print("Pooling Image Shape:", max_pooling_img.shape)
```

```
Input Image Shape: torch.Size([1, 1, 225, 225])
```

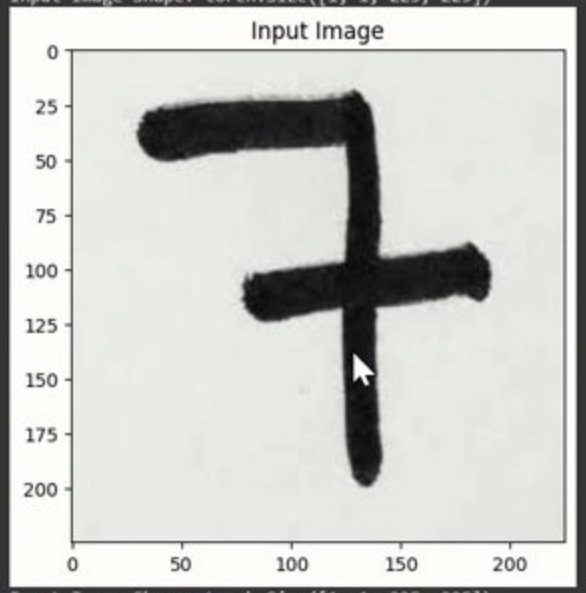


```
Input Image Shape: torch.Size([1, 1, 225, 225])
```

Conv Image

1s completed at 9:43 PM

AISCIENCES

Pytorch.ipynb - Colab | CNN Explainer | Animations of Convolution and | Image Kernels explained visuall

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=sud58LClW23n

Paused

+ Code + Text Saving...

RAM Disk

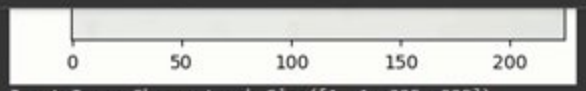


```
Input Image Shape: torch.Size([1, 1, 225, 225])
```

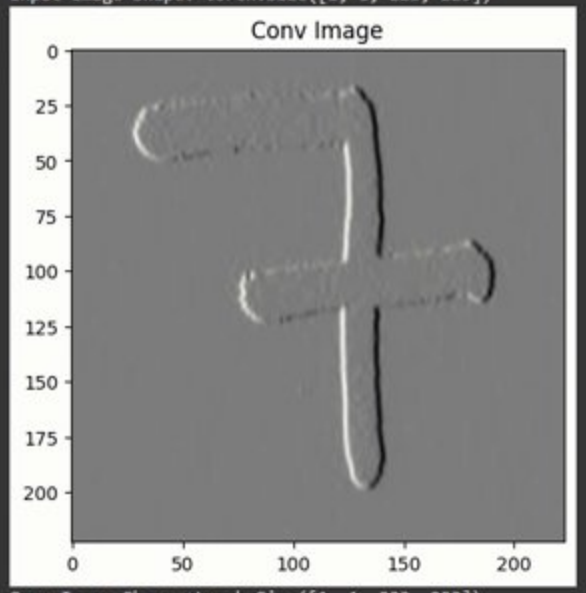


```
Conv Image Shape: torch.Size([1, 1, 223, 223])
```

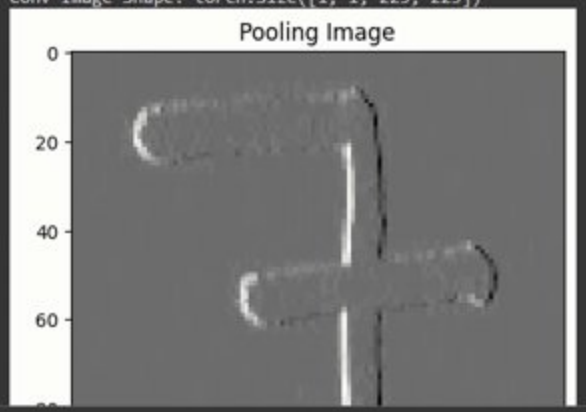

1s completed at 9:43 PM

AISCIENCES

Pytorch.ipynb - Colab | CNN Explainer | Animations of Convolution and | Image Kernels explained visuall

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=sud58LClW23n

Paused

+ Code + Text All changes saved

RAM Disk

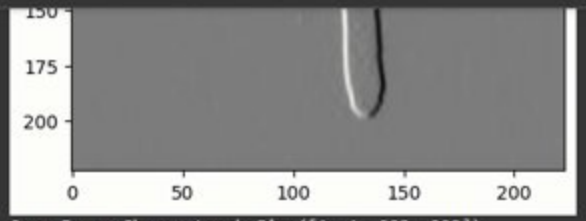

Conv Image Shape: torch.Size([1, 1, 223, 223])

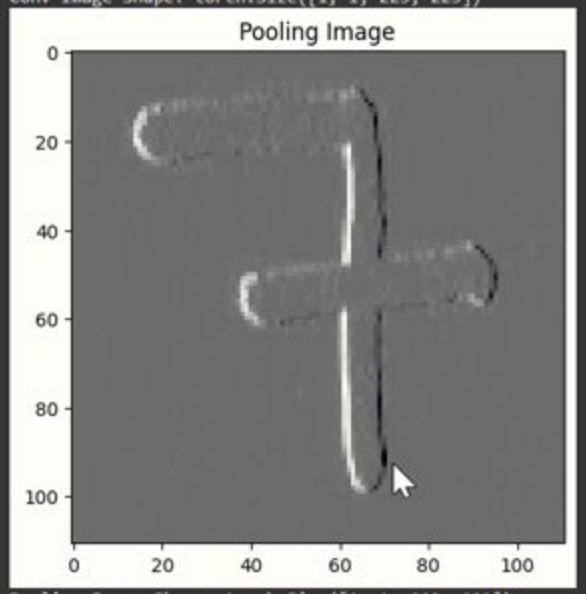

Pooling Image Shape: torch.Size([1, 1, 111, 111])

[ ]

[ ]

[ ]

[ ]

1s completed at 9:43 PM

AISCIENCES

Pytorch.ipynb - Colab | CNN Explainer | Animations of Convolution and | Image Kernels explained visuall

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=sud58LClW23n

Paused

+ Code + Text All changes saved

RAM Disk

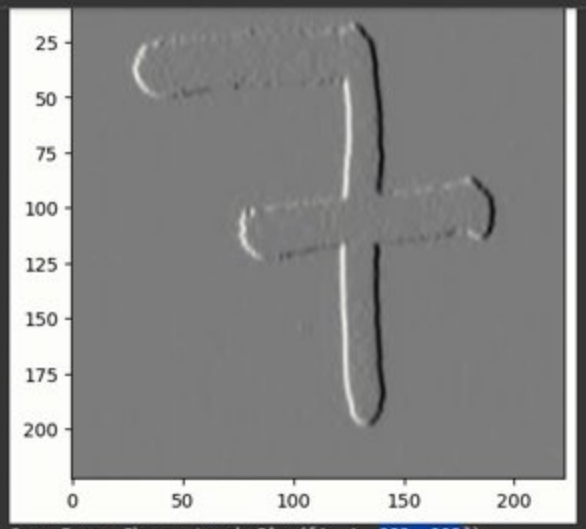


```
Conv Image Shape: torch.Size([1, 1, 223, 223])
```

[ ]

[ ]

[ ]

[ ]

1s completed at 9:43 PM

AISCIENCES

Pytorch.ipynb - Colab | CNN Explainer | Animations of Convolution and | Image Kernels explained visuall

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=sud58LClW23n

Paused

+ Code + Text All changes saved

RAM Disk

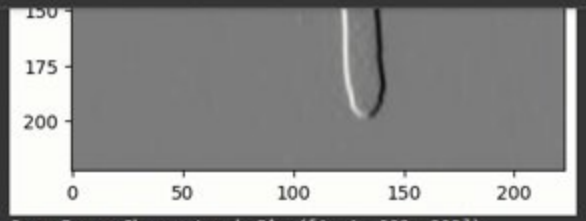

Conv Image Shape: torch.Size([1, 1, 223, 223])

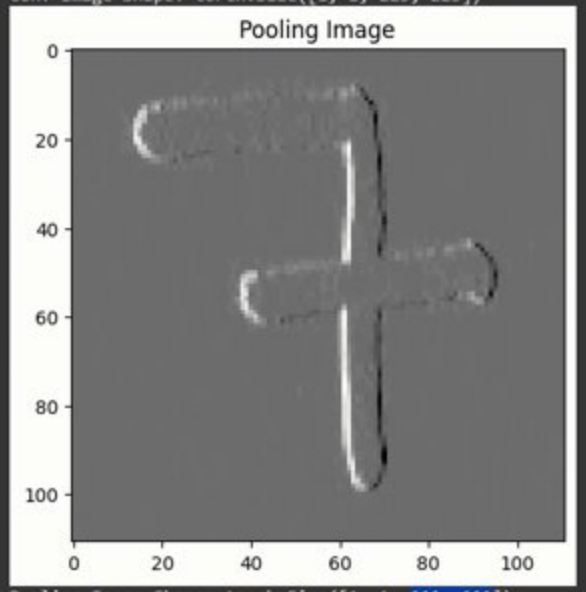

Pooling Image Shape: torch.Size([1, 1, 111, 111])

[ ]

[ ]

[ ]

[ ]

1s completed at 9:43 PM

AISCIENCES

Pytorch.ipynb - Colab | CNN Explainer | Animations of Convolution and | Image Kernels explained visuall

poloclub.github.io/cnn-explainer/

Paused

CNN EXPLAINER Learn Convolutional Neural Network (CNN) in your browser!

Show detail | Unit

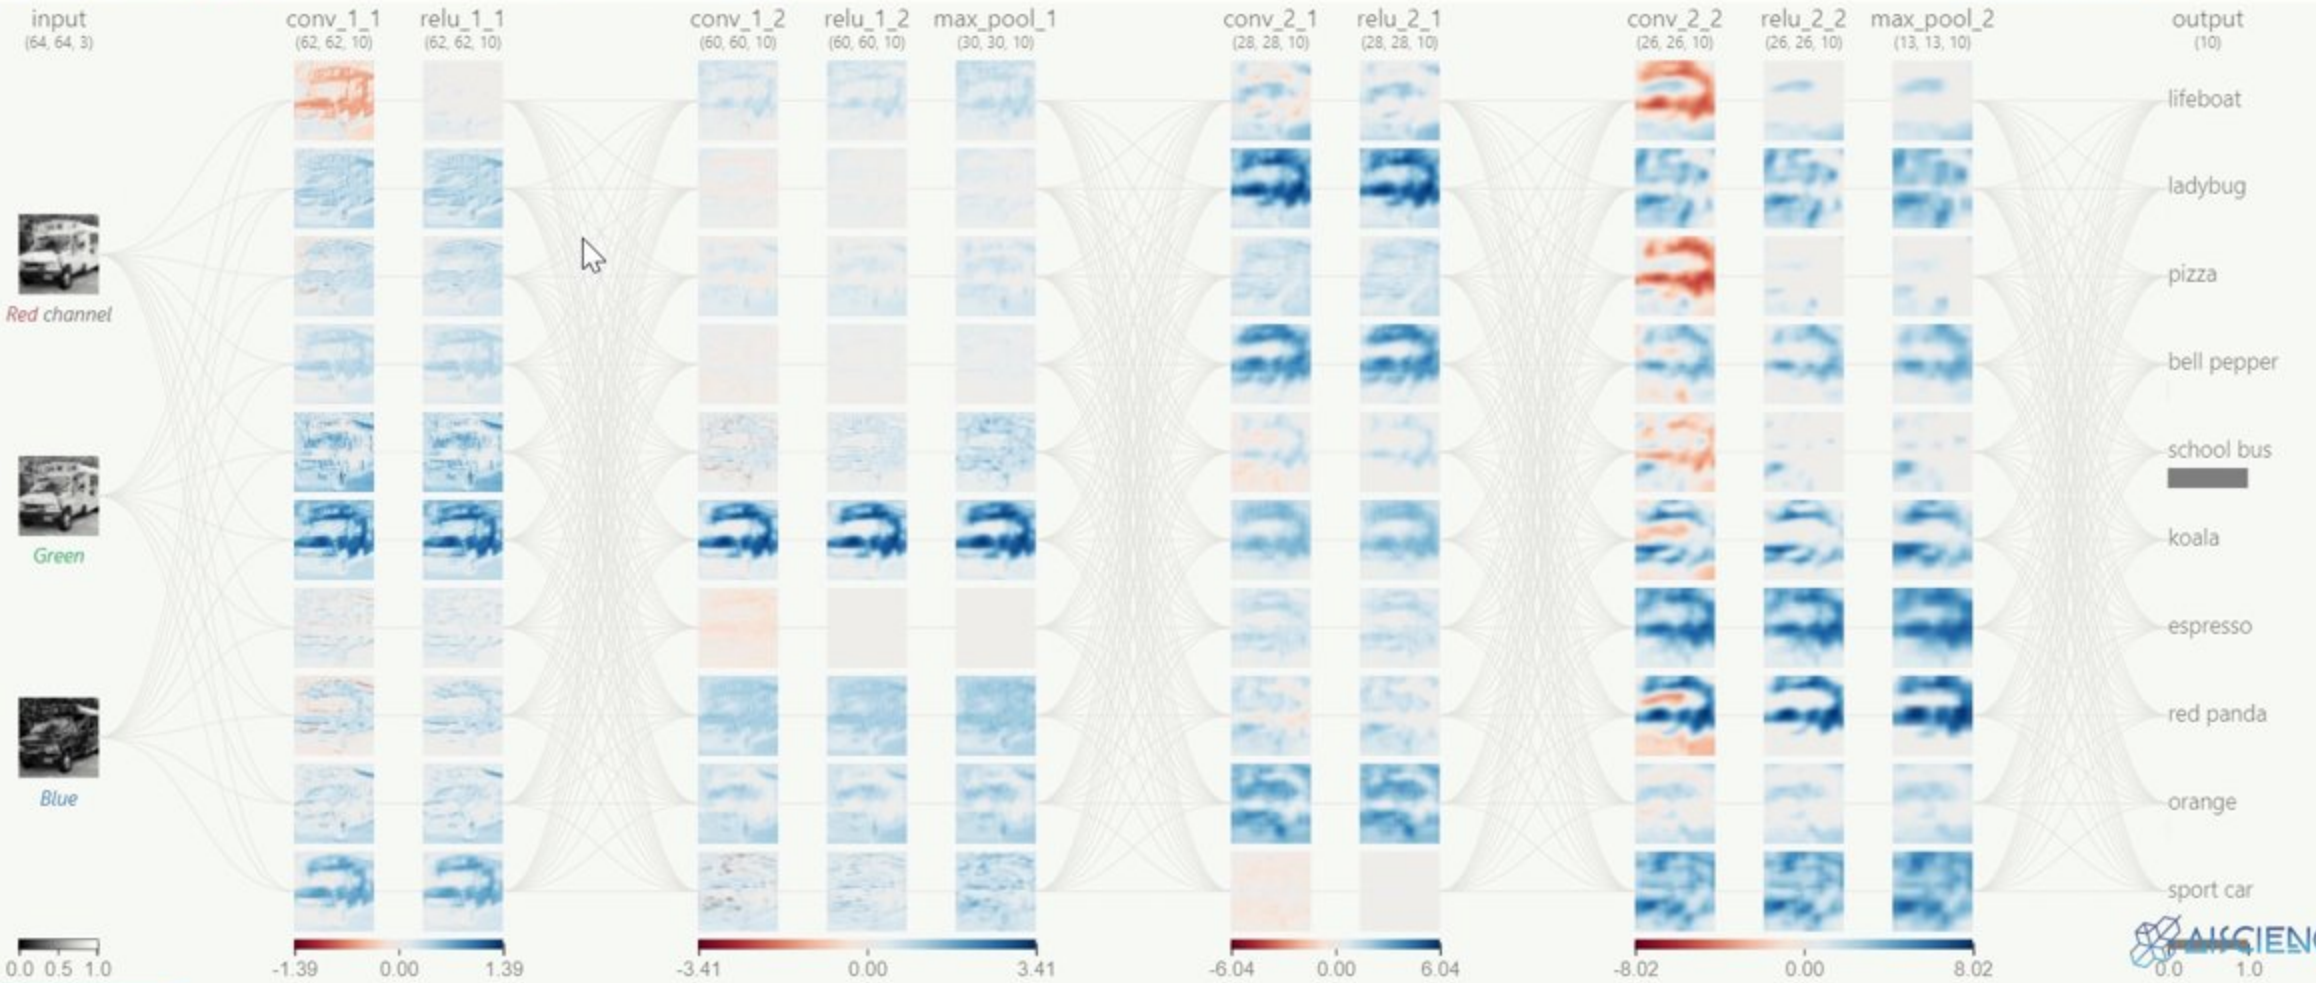

Pytorch.ipynb - Colab | CNN Explainer | Animations of Convolution and | Image Kernels explained visuall

poloclub.github.io/cnn-explainer/

Paused

CNN EXPLAINER Learn Convolutional Neural Network (CNN) in your browser!

Show detail | Unit

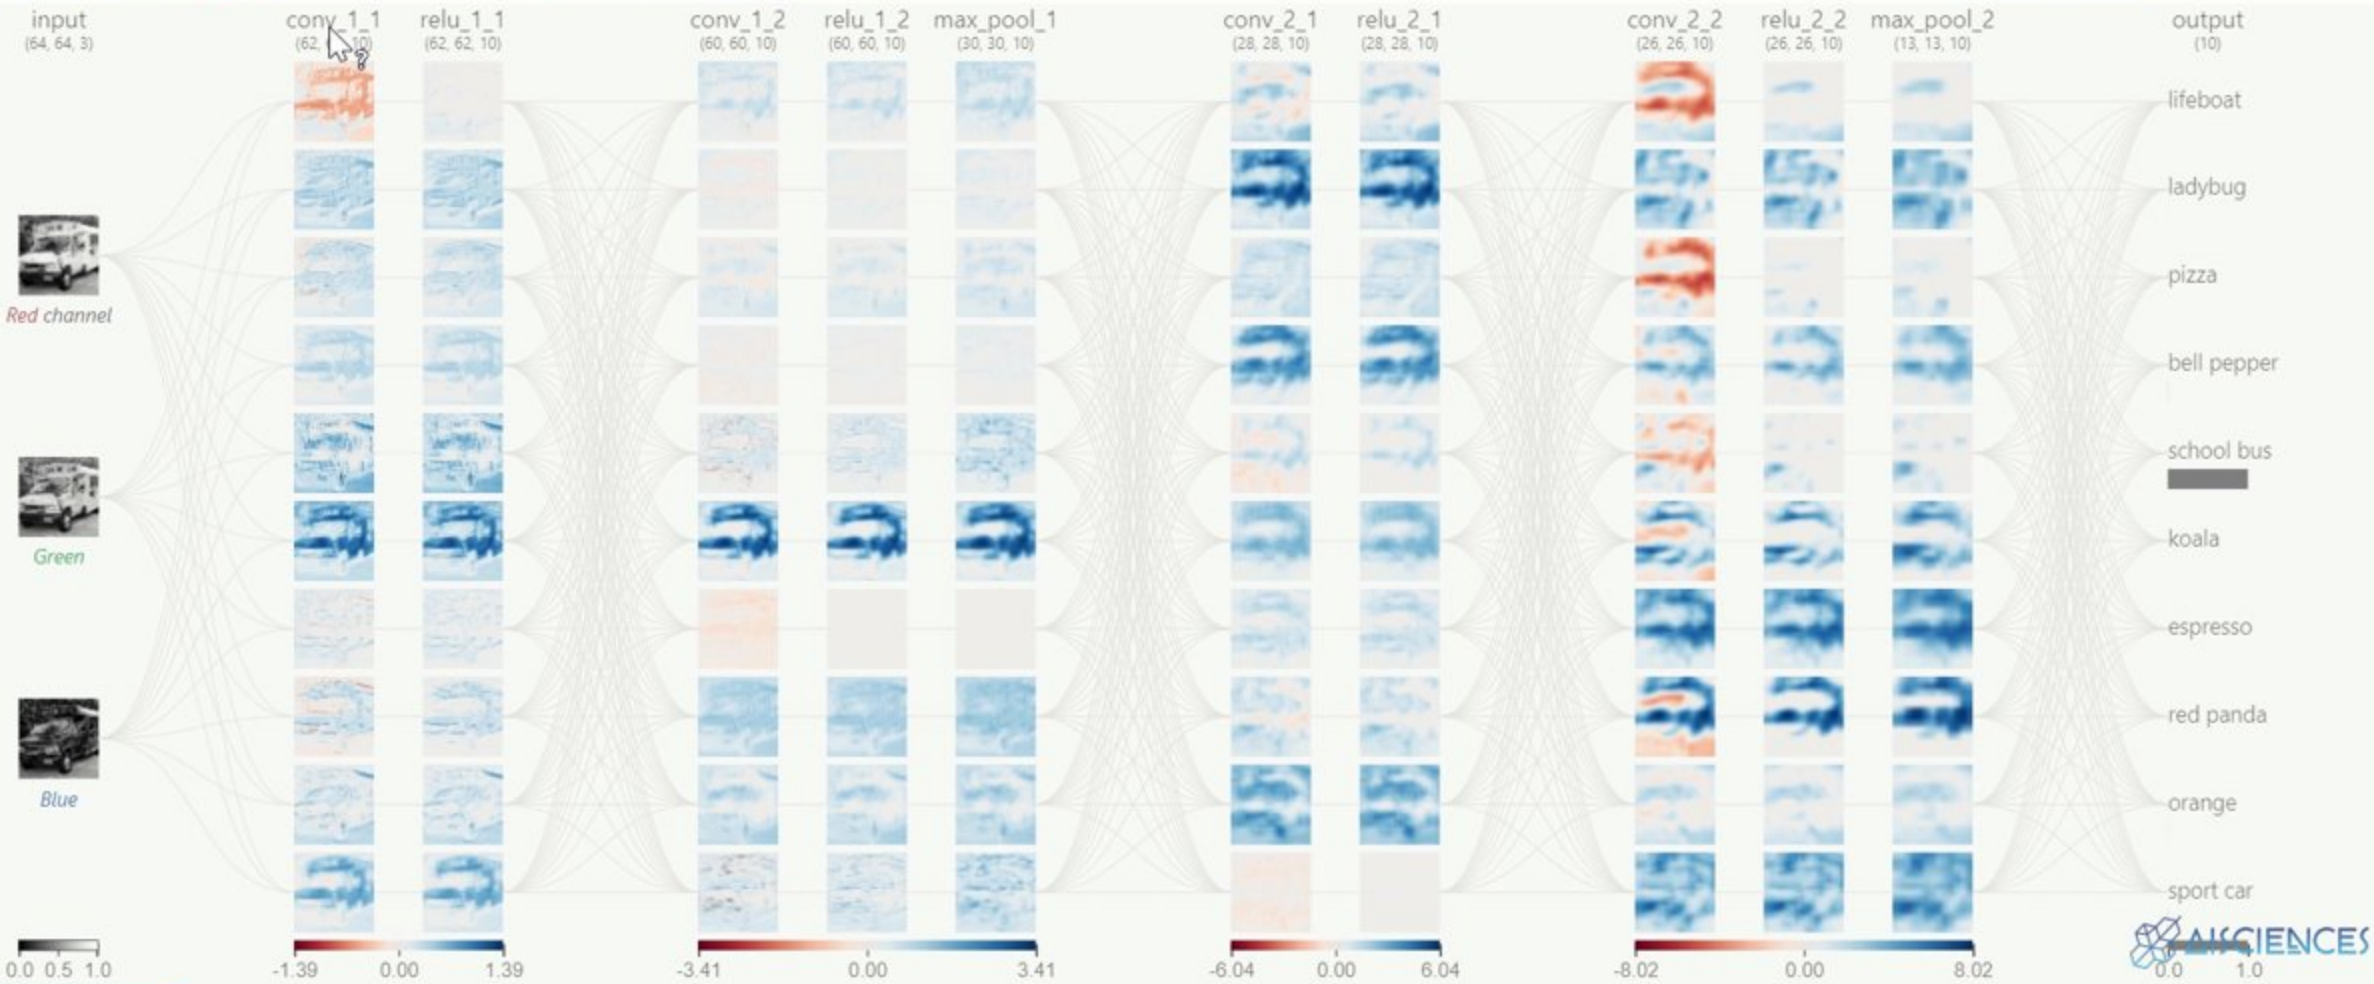

Pytorch.ipynb - Colab | CNN Explainer | Animations of Convolution and | Image Kernels explained visuall

poloclub.github.io/cnn-explainer/

Paused

CNN EXPLAINER Learn Convolutional Neural Network (CNN) in your browser!

Show detail | Unit

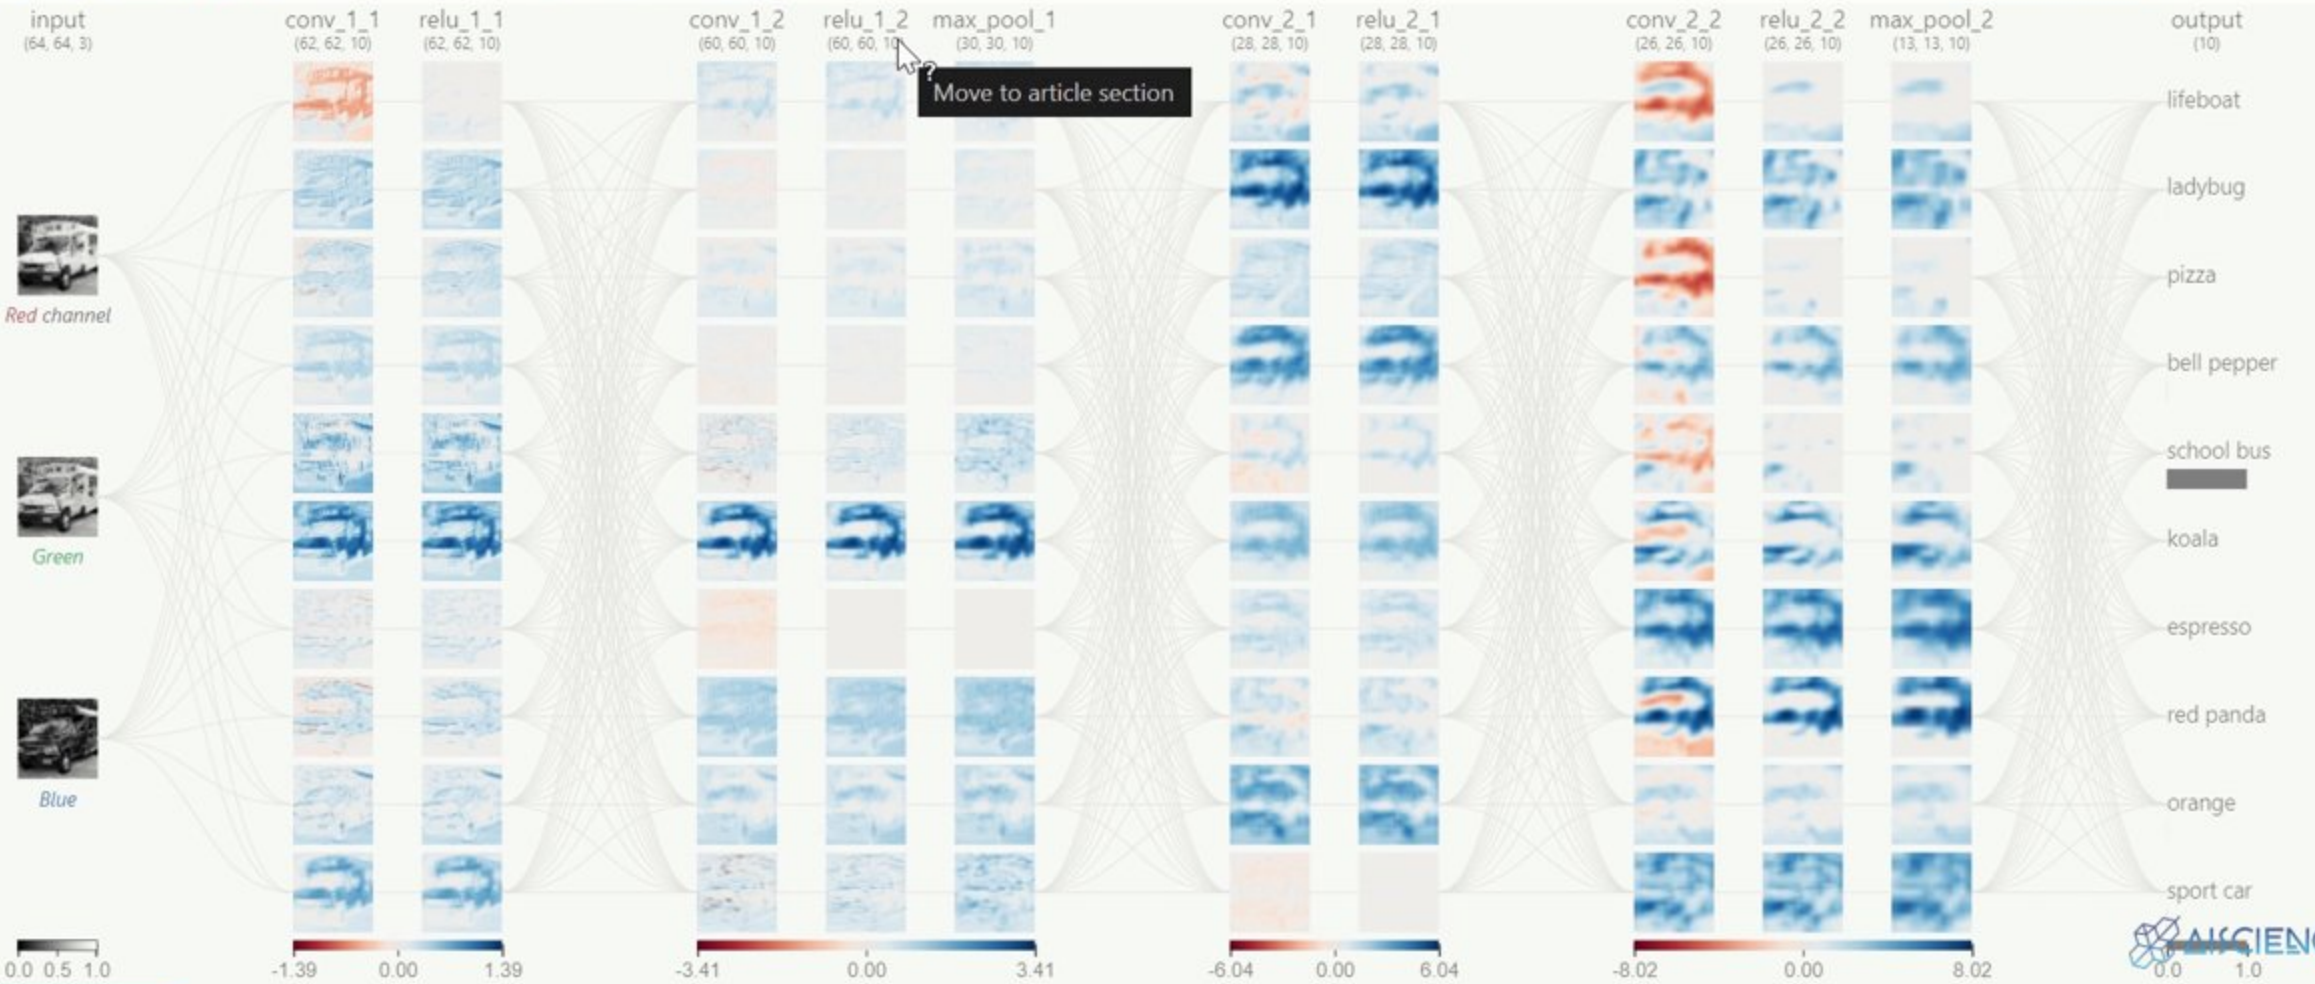

Pytorch.ipynb - Colab | CNN Explainer | Animations of Convolution and | Image Kernels explained visuall

poloclub.github.io/cnn-explainer/

Paused

CNN EXPLAINER Learn Convolutional Neural Network (CNN) in your browser!

Show detail | Unit

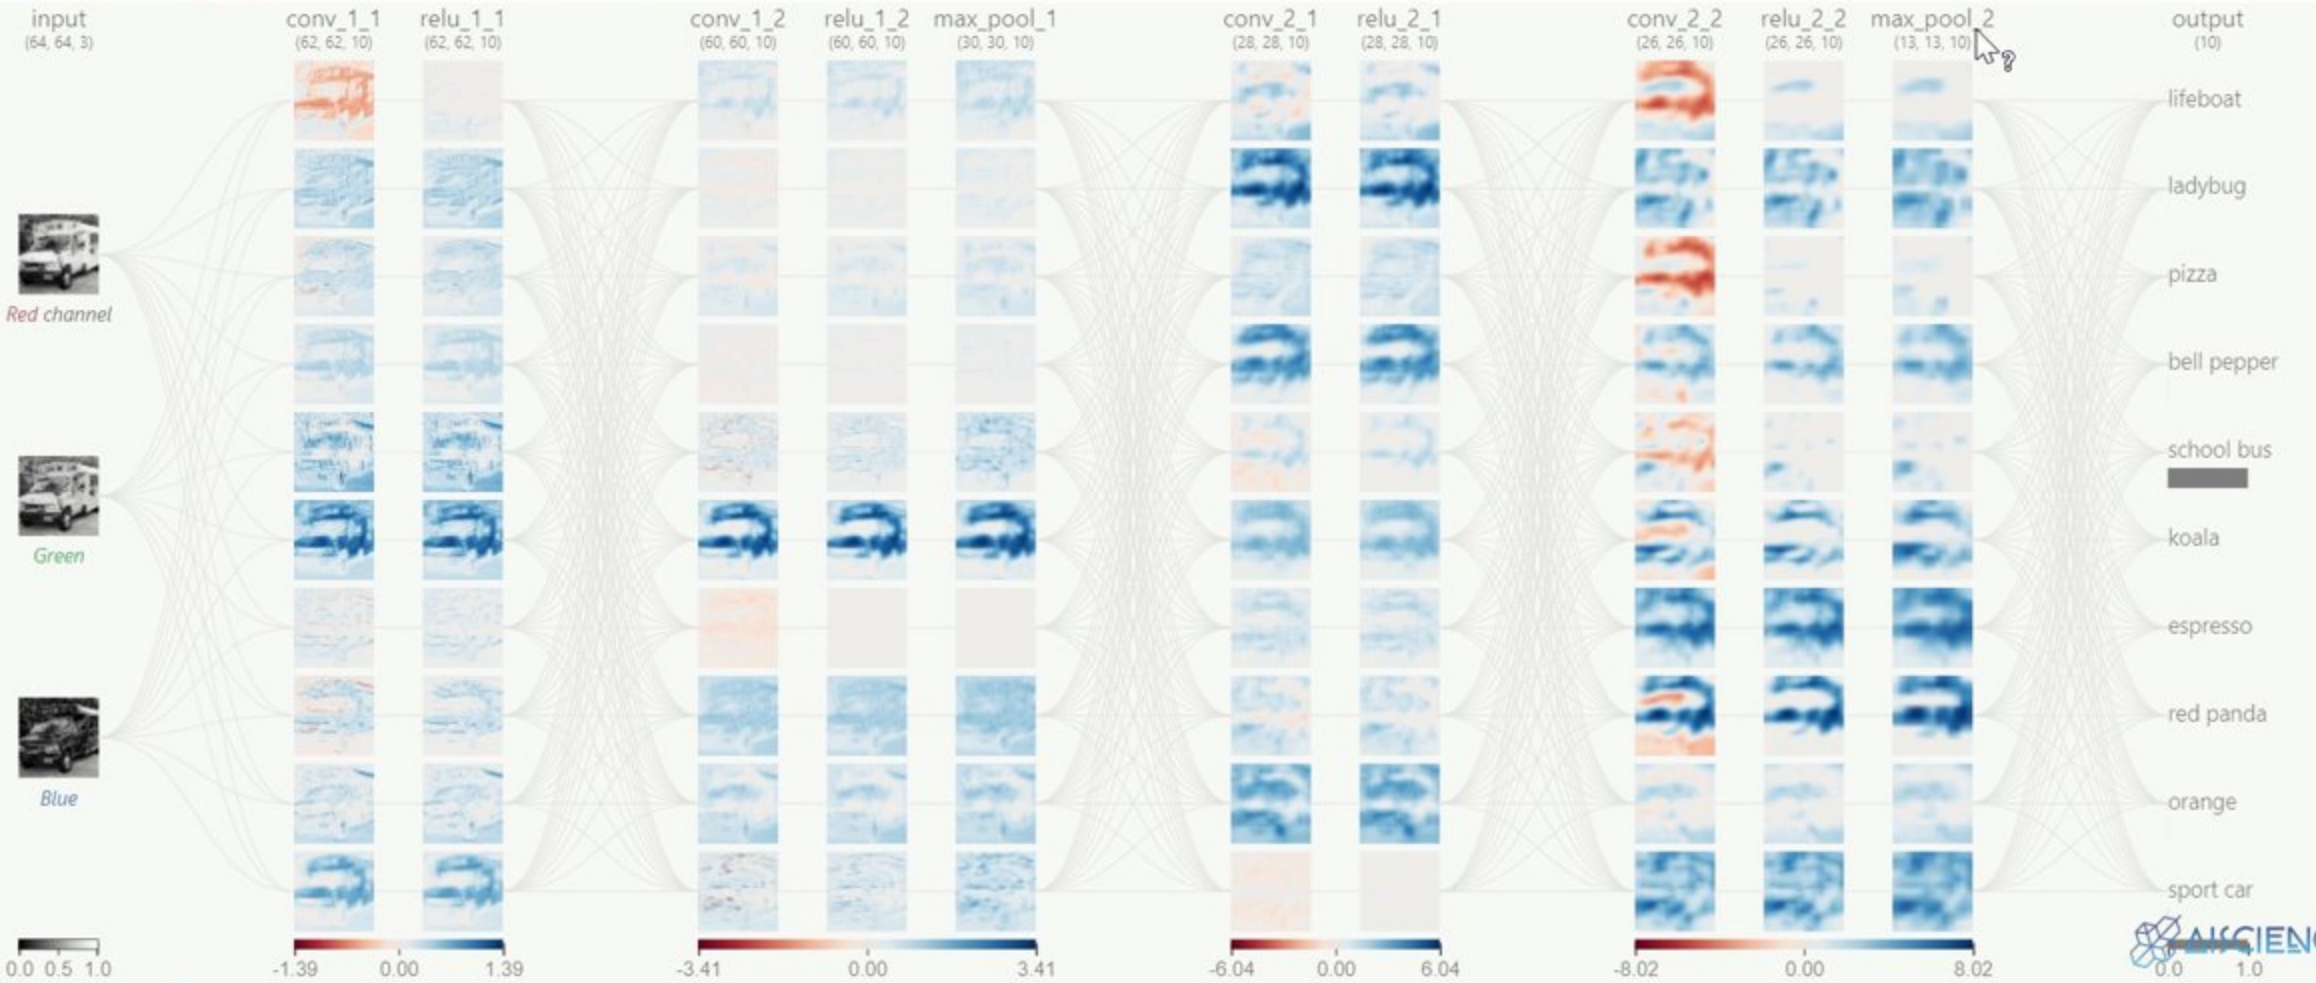

Pytorch.ipynb - Colab | CNN Explainer | Animations of Convolution and | Image Kernels explained visuall

poloclub.github.io/cnn-explainer/

Paused

CNN EXPLAINER Learn Convolutional Neural Network (CNN) in your browser!

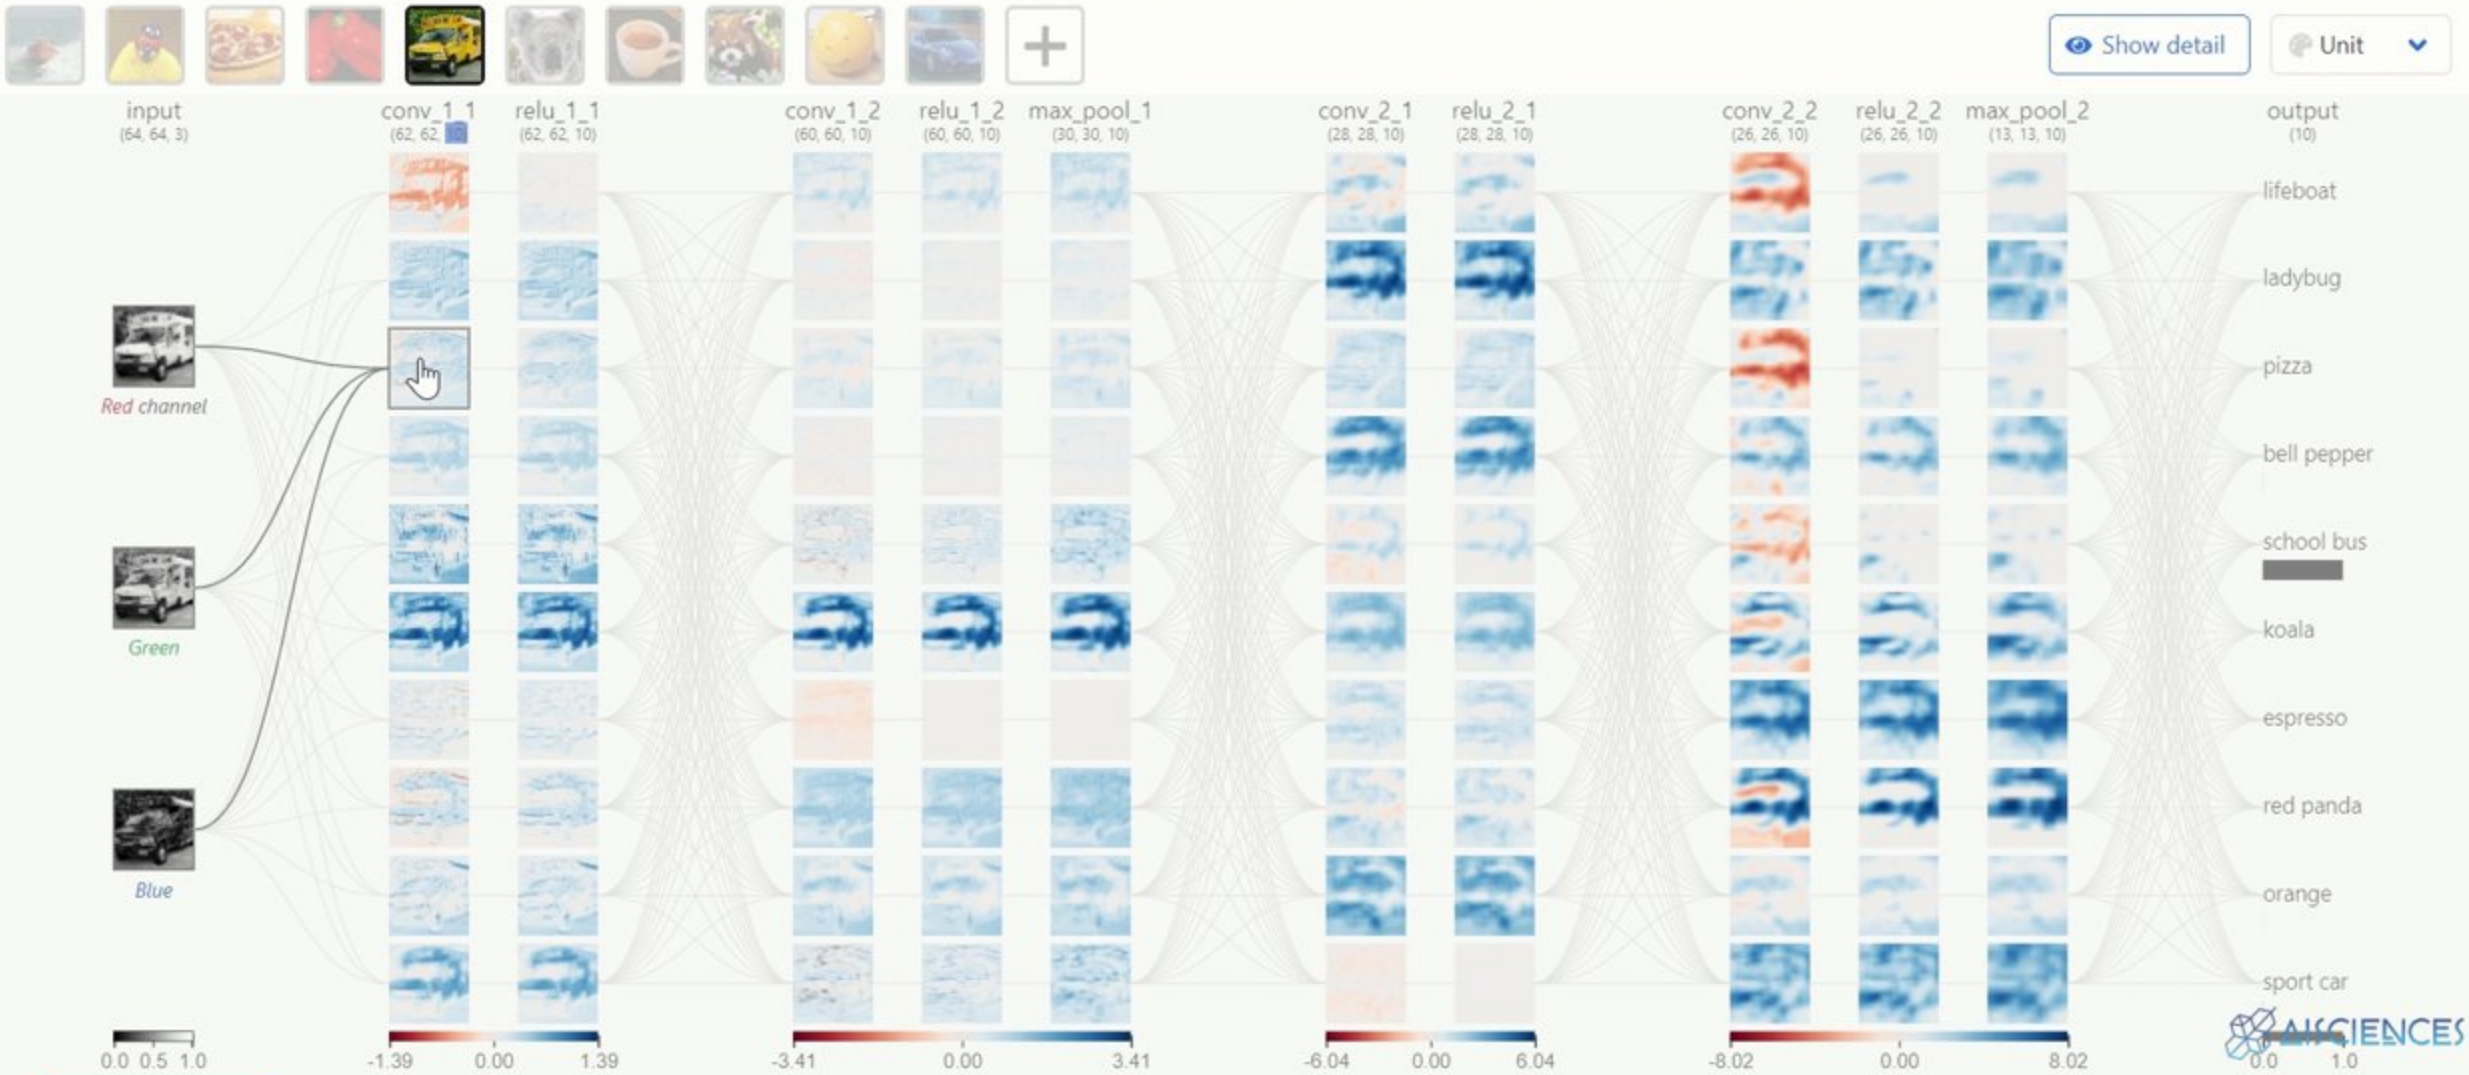

Pytorch.ipynb - Colab | CNN Explainer | Animations of Convolution and | Image Kernels explained visuall

poloclub.github.io/cnn-explainer/

Paused

CNN EXPLAINER Learn Convolutional Neural Network (CNN) in your browser!

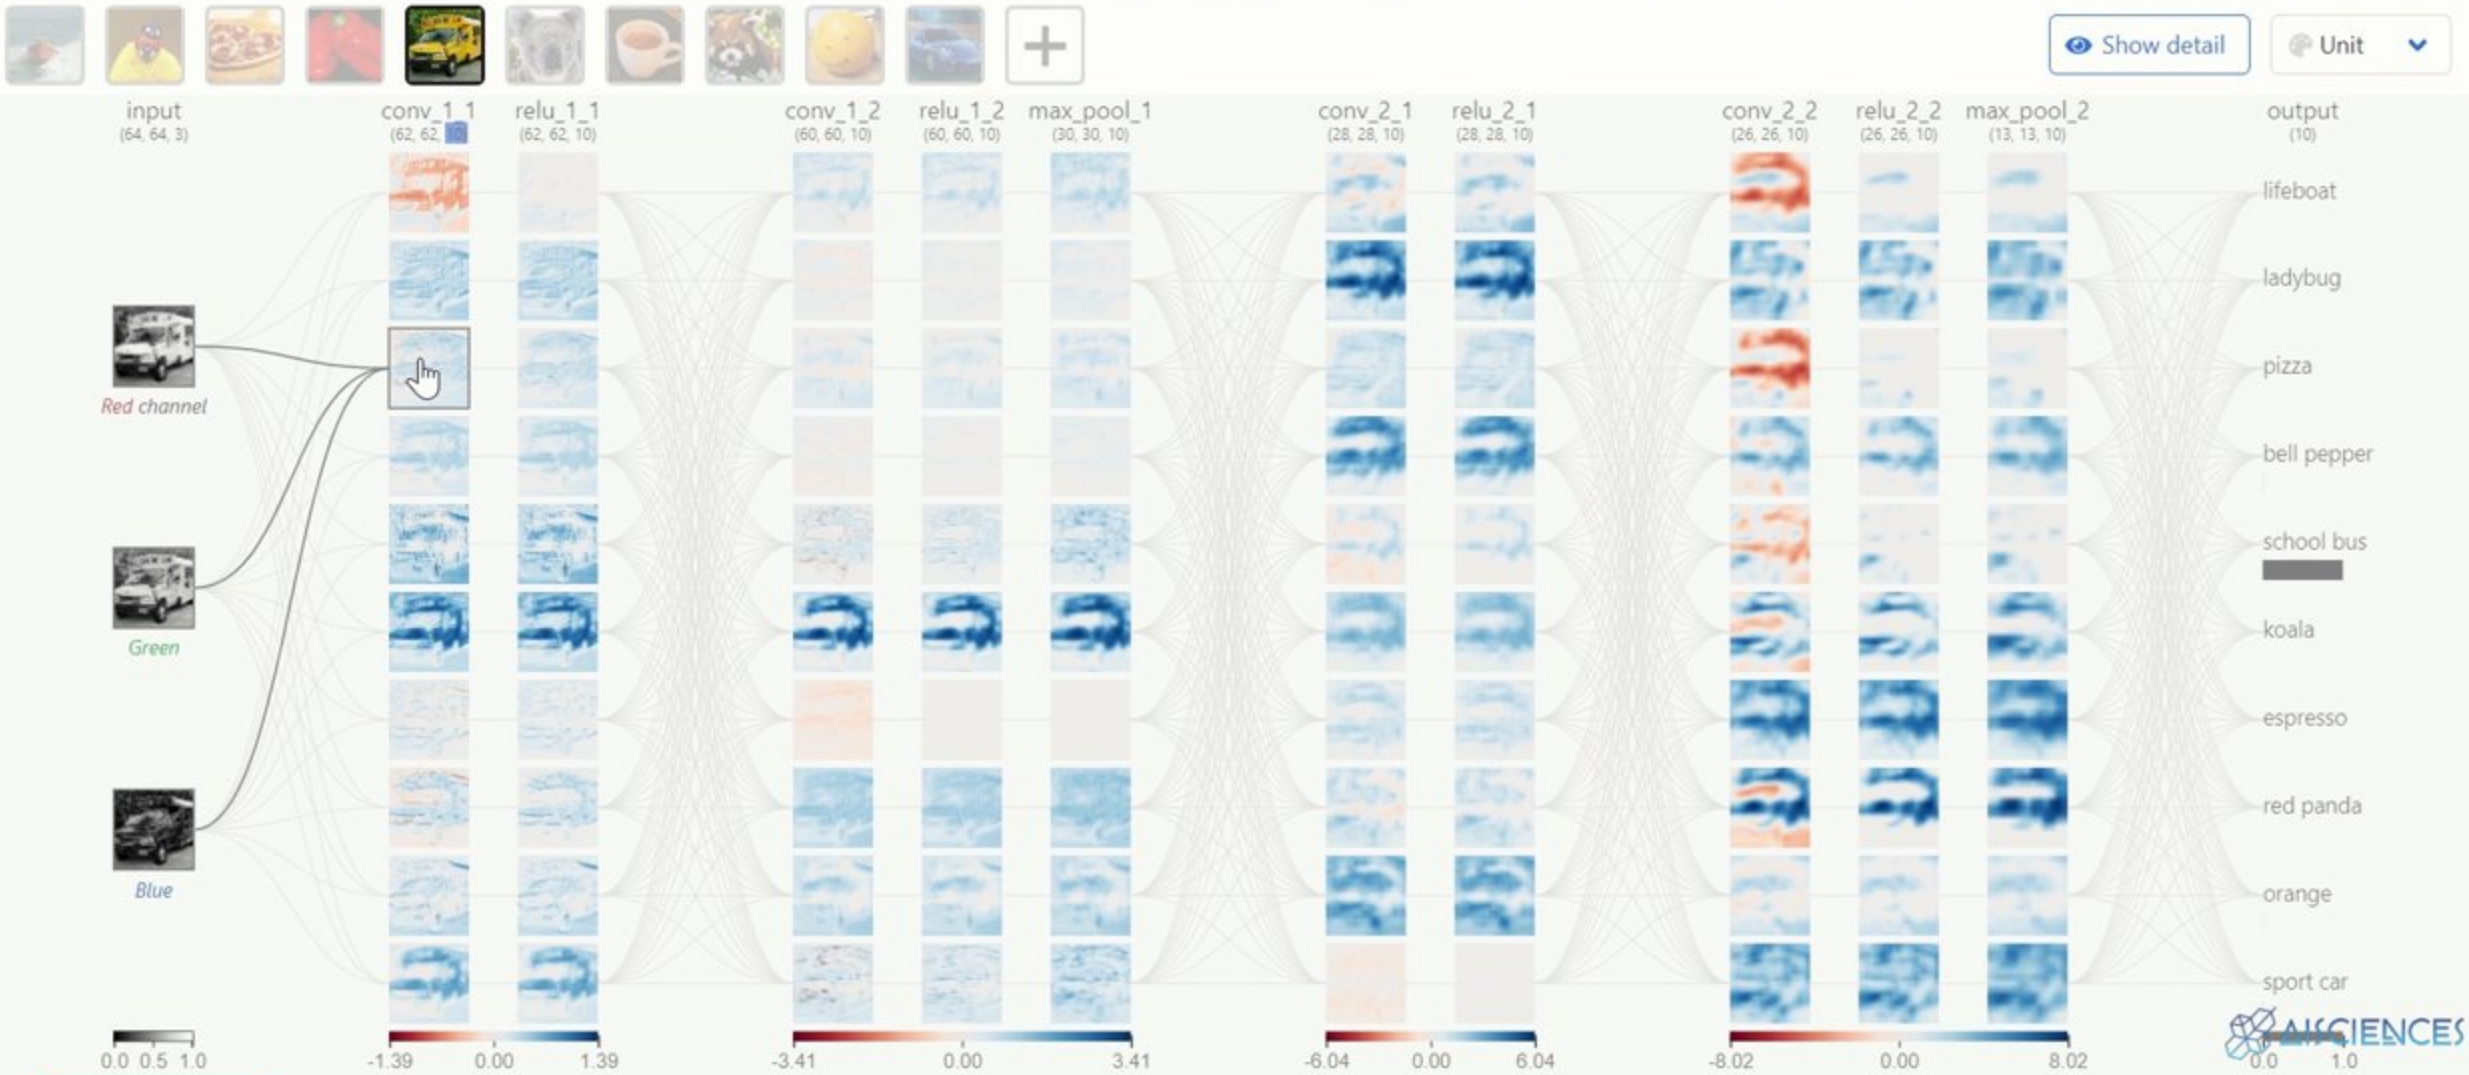

Pytorch.ipynb - Colab | CNN Explainer | Animations of Convolution and | Image Kernels explained visuall

poloclub.github.io/cnn-explainer/

Paused

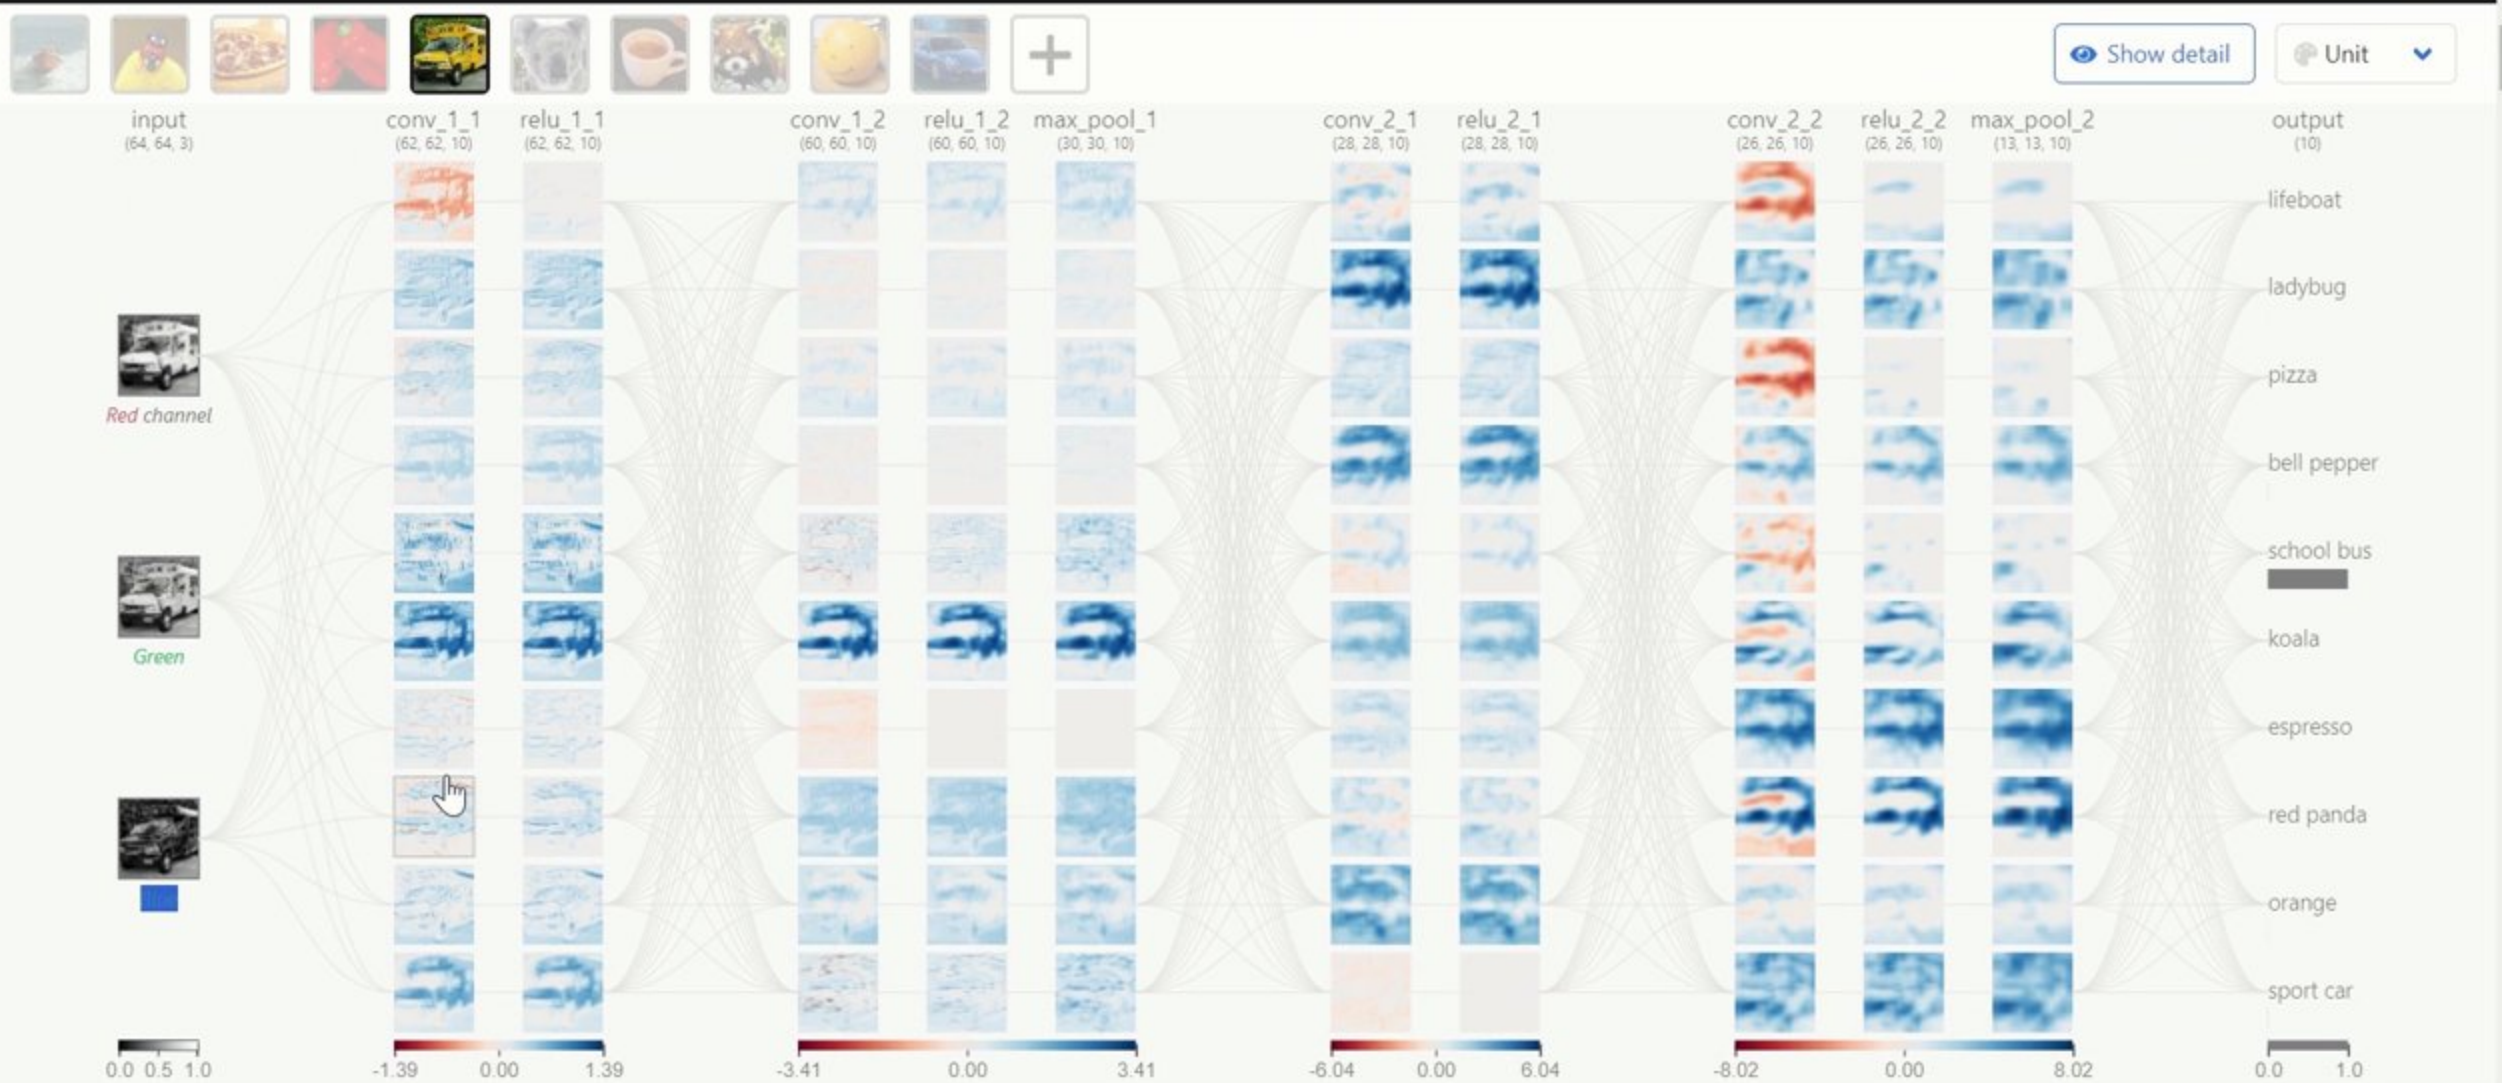

AISCIENCES

Pytorch.ipynb - Colab | CNN Explainer | Animations of Convolution and | Image Kernels explained visuall

poloclub.github.io/cnn-explainer/

Paused

AISCIENCES

Pytorch.ipynb - Colab | CNN Explainer | Animations of Convolution and | Image Kernels explained visuall

poloclub.github.io/cnn-explainer/

Paused

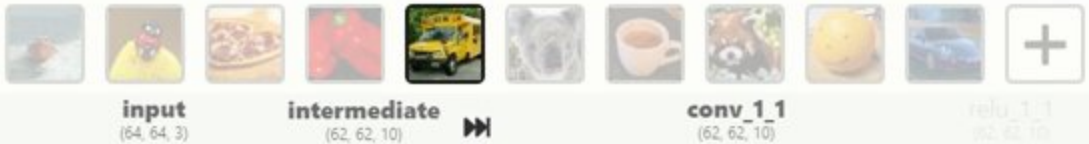

input
(64, 64, 3)

intermediate
(62, 62, 10)

conv_1_1
(62, 62, 10)

Show detail

Unit

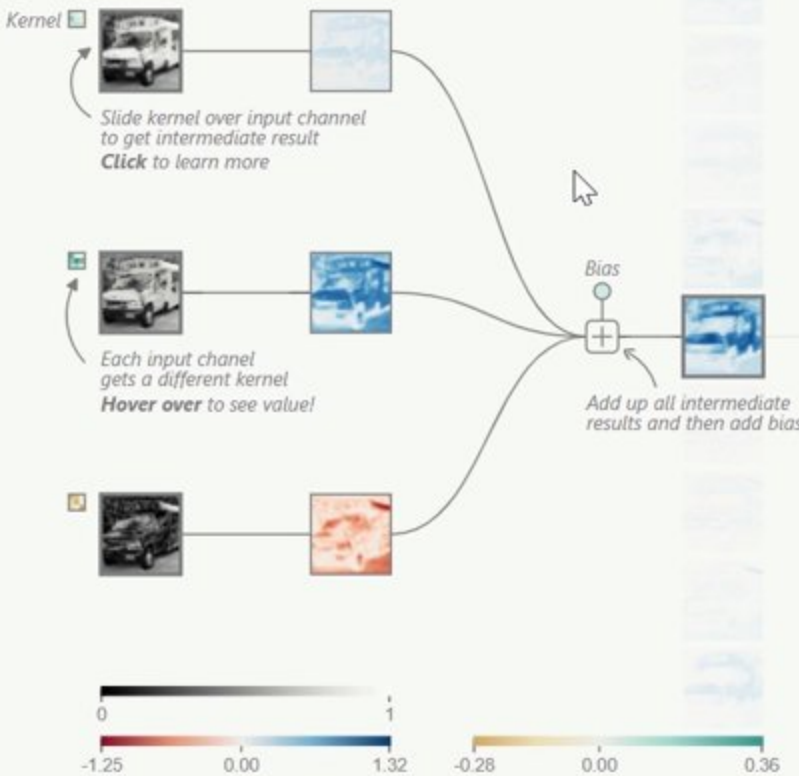

AISCIENCES

Udemy

Pytorch.ipynb - Colab | CNN Explainer | Animations of Convolution and | Image Kernels explained visuall

poloclub.github.io/cnn-explainer/

Paused

Show detail | Unit

relu_1_2 (60, 60, 10) | max_pool_1 (30, 30, 10)

-3.41 0.00 3.41

Max Pooling

Input (60, 60)

max(1.79 1.96 / 1.71 1.72) = 1.96

Output (30, 30)

Hover over the matrices to change kernel position.

AISCIENCES

Udemy

Pytorch.ipynb - Colab | CNN Explainer | Animations of Convolution and | Image Kernels explained visuall

poloclub.github.io/cnn-explainer/

Paused

Show detail | Unit

relu_1_2 (60, 60, 10) | max_pool_1 (30, 30, 10)

-3.41 0.00 3.41

Max Pooling

Input (60, 60)

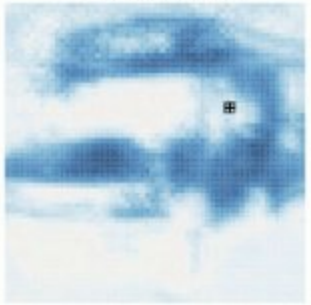

max(0.9 0.85 / 0.79 0.75) = 0.9

Output (30, 30)

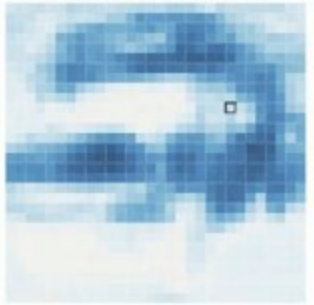

Hover over the matrices to change kernel position.

AISCIENCES

Udemy

Pytorch.ipynb - Colab | CNN Explainer | Animations of Convolution and | Image Kernels explained visuall

poloclub.github.io/cnn-explainer/#article-pooling

Paused

CNN EXPLAINER Learn Convolutional Neural Network (CNN) in your browser!

Show detail | Unit

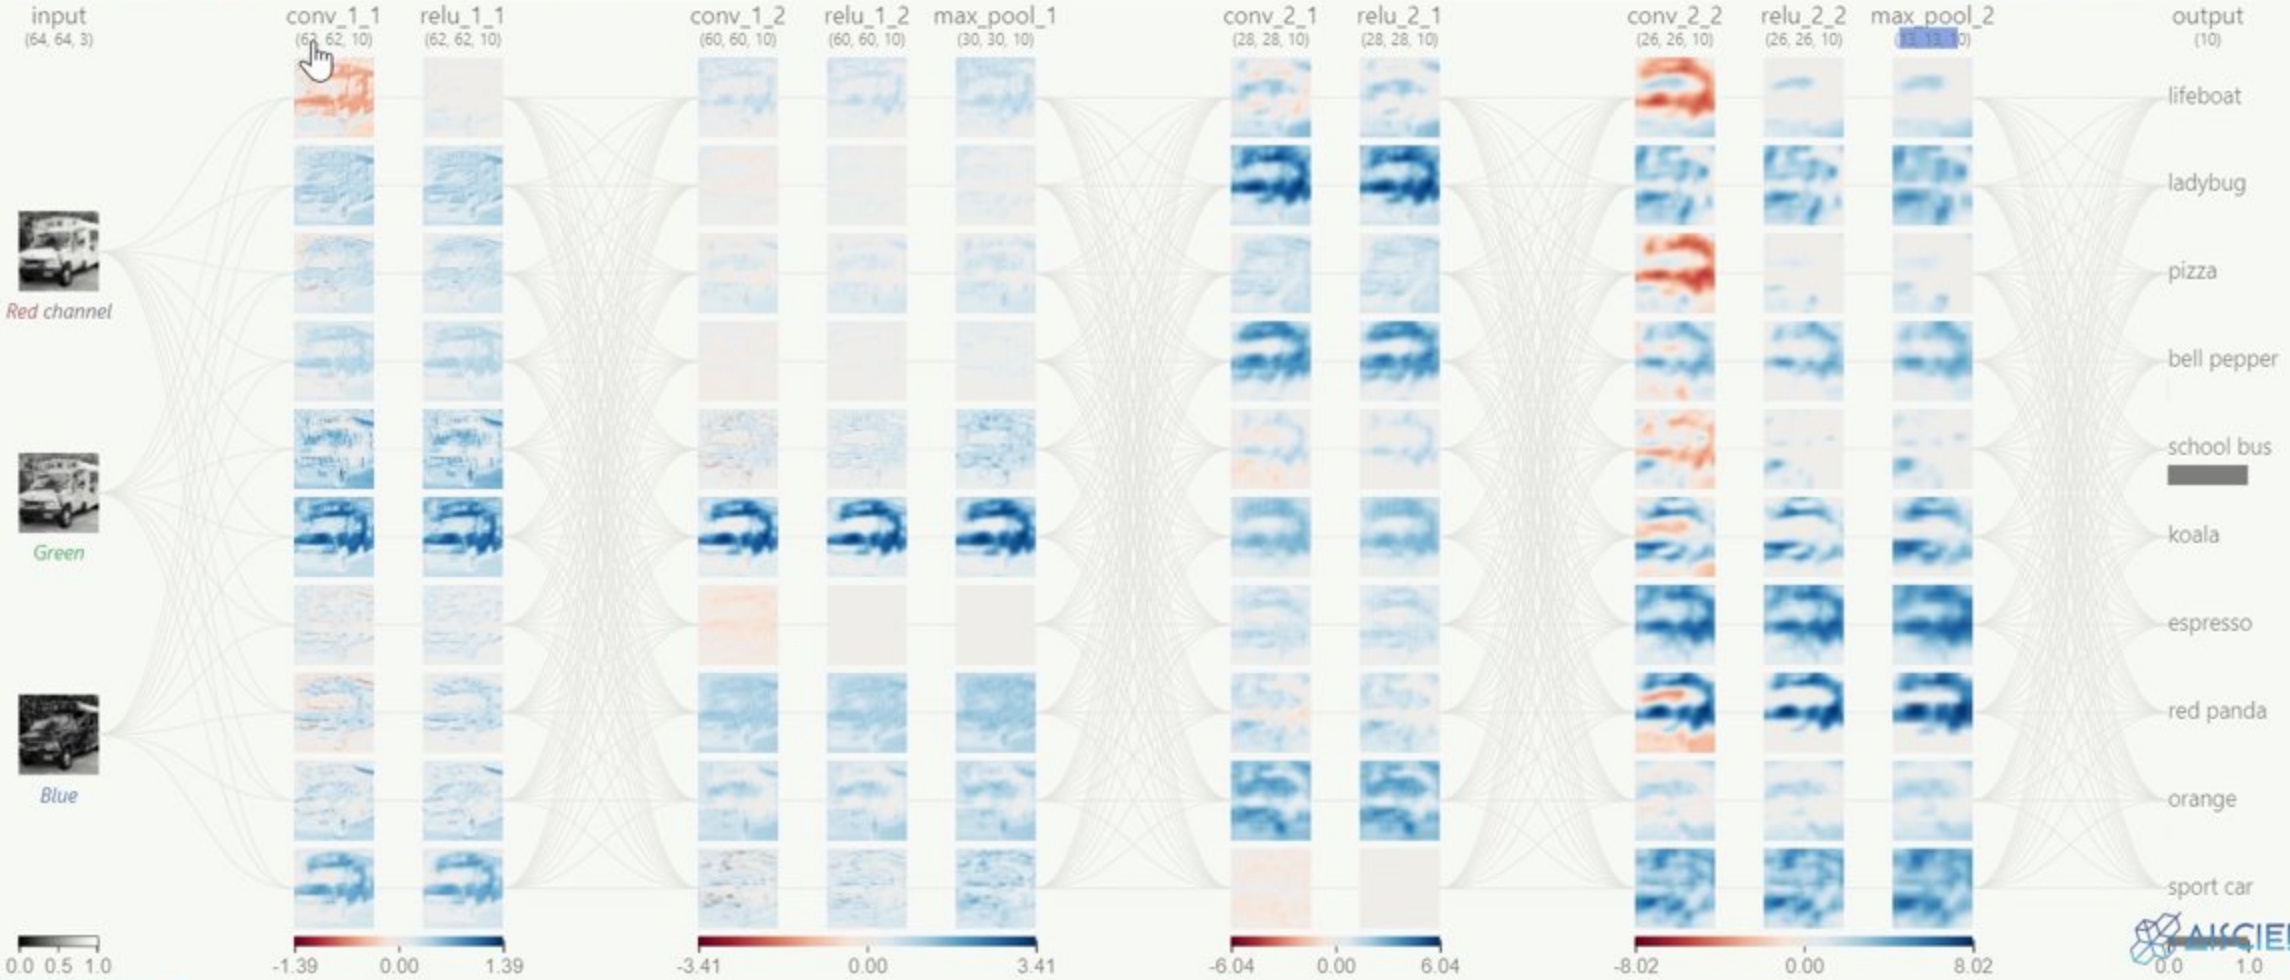

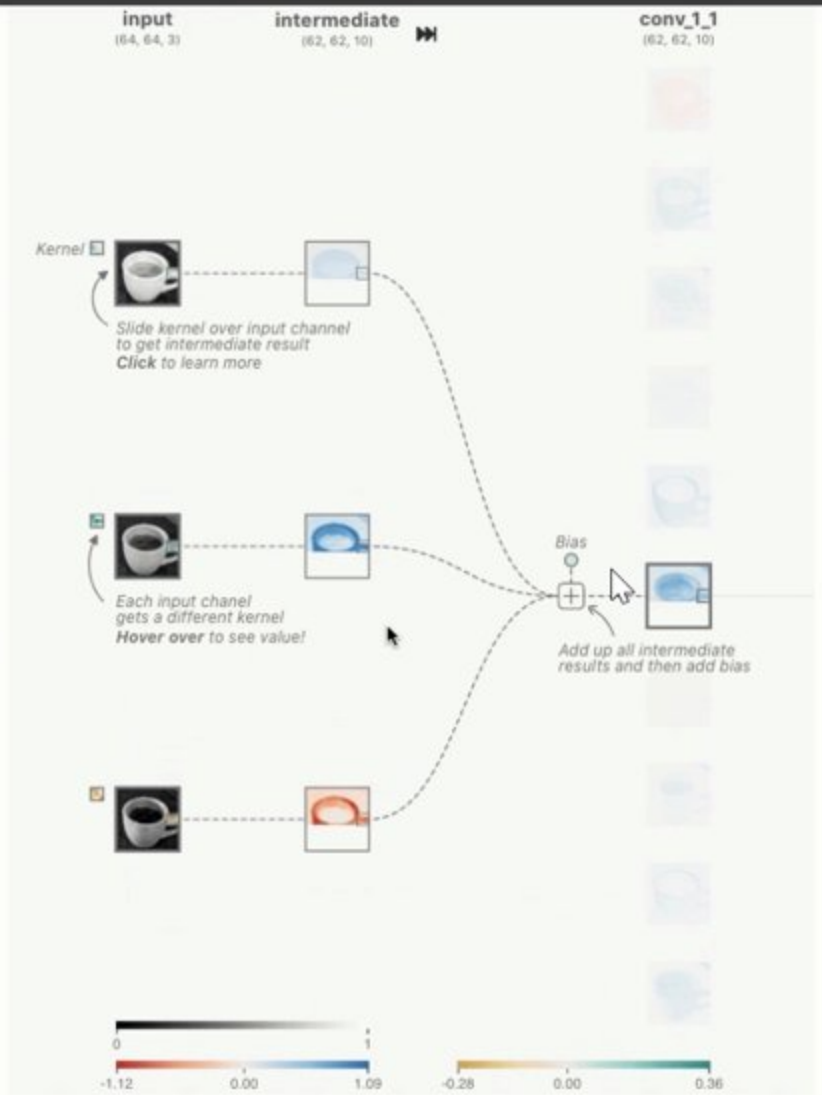

Figure 1. As you hover over the activation map of the topmost node from the first convolutional layer, you can see that 3 kernels were applied to yield this activation map. After clicking this activation map, you can see the

MAX POOLING

Input

$$\begin{bmatrix} 6 & 3 & 2 \\ 5 & 2 & 1 \\ 4 & 1 & 0 \end{bmatrix}$$

Pooling matrix 2x2

Output

$$\begin{bmatrix} 6 & 2 \\ 4 & 0 \end{bmatrix}$$

CONVOLUTION

Input

$$\begin{bmatrix} 1 & 2 & 0 & 1 & 3 \\ 4 & 1 & 0 & 2 & 1 \\ 2 & 3 & 1 & 0 & 1 \\ 1 & 2 & 0 & 1 & 3 \\ 3 & 1 & 2 & 3 & 0 \end{bmatrix}$$

Kernal

$$\begin{bmatrix} 1 & 0 & -1 \\ 1 & 0 & -1 \\ 1 & 0 & -1 \end{bmatrix}$$

Output

$$\begin{bmatrix} 6 & 3 & 2 \\ 5 & 2 & 1 \\ 4 & 1 & 0 \end{bmatrix}$$

$$\text{Output}(1,2) = (2\times1)+(0\times0)+(1\times-1)+(1\times1)+(0\times0)+(2\times-1)+(3\times1)+(1\times0)+(0\times-1)$$

$$=2+0-1+1+0-2+3+0-0$$

$$=3$$

CODES & DATA ARE AVAILBLE AT
WWW.AISCIENCES.ACADEMY/COURSE-PYTORCH