


# IMPORTANT ELEMENTS OF NEURAL NETWORK

- ← • Loss Function
- Optimizer
- Activation Function




# IMPORTANT ELEMENTS OF NEURAL NETWORK

- Loss Function
- Optimizer
- Activation Function



# SIMPLE NEURAL NETWORK



$$X * W + b$$





Pytorch.ipynb - Colab

Loss Functions.pdf

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=6DG1zS7iKBCP

✓ T4 RAM Disk

Paused

+ Code + Text All changes saved

✓ 2s

[2]

tensor([[ 105., 120., 75.],  
[ 78., 90., 50.],  
[ 82., 70., 45.]])

✓ 0s

w = torch.randn(3,2, requires\_grad=True)  
b = torch.randn(2, requires\_grad=True)  
print(w)  
print(b)

tensor([[ 0.1674, 0.2380],  
[ 0.8814, -0.4406],  
[ 0.9254, 0.4304]], requires\_grad=True)  
tensor([ 1.2801, -0.7225], requires\_grad=True)

✓ 0s

[9] # Defining Model  
def model(x):  
# return (x @ w) + b  
return torch.matmul(x,w)+b  
  
#MSE Loss function

↑ ↓ ↶ ↷ ⚙ 📄 🗑 ⋮

AI SCIENCES

Connected to Python 3 Google Compute Engine backend (GPU)

Udacity

I

```
preds = model(inputs)
loss = mse(preds, targets)
print(loss)
```

Pytorch.ipynb - Colab

Loss Functions.pdf

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=Oor1qUGUqzSO

+

Code

+

Text

All changes saved

✓

0s

[9]

167.897216796875

97.26914978027344

63.02490234375

46.403343200683594

38.319190979003906

34.37253952026367

32.43240737915039

31.466655731201172

30.975189208984375

30.71561050415039

tensor(30.7137, grad\_fn=<DivBackward0>)

[ ]

[ ]

↑

↓

↻

💬

✎

📄

🗑️

⋮

>

Basic Neural Network with Built-in functions

Using Pytorch

🎮

↳ 37 cells hidden

ASCIENCES

✓

Connected to Python 3 Google Compute Engine backend (GPU)

Udemy

Pytorch.ipynb - Colab

Loss Functions.pdf

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=Oor1qUGUqzSO

☆

Paused

+ Code + Text All changes saved

T4

RAM

Disk

import torch

import numpy as np

# Input (temp, rainfall, humidity)

inputs = np.array([[73, 67, 43],

[91, 88, 64],

[87, 134, 58],

[102, 43, 37],

[69, 96, 70],

[85, 100, 60],

[95, 80, 55],

[105, 120, 75],

[78, 90, 50],

[82, 70, 45]], dtype='float32')

# Targets (apples, oranges)

targets = np.array([[56, 70],

[81, 101],

[119, 133],

[22, 37],

[103, 119],

[98, 110],

[88, 95],

[115, 140],

[76, 85],

[65, 75]], dtype='float32')

inputs = torch.from\_numpy(inputs)

targets = torch.from\_numpy(targets)

print(inputs)

print(targets)

Connected to Python 3 Google Compute Engine backend (GPU)



Pytorch.ipynb - Colab

Loss Functions.pdf

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=B6Ymnu54qzSb

+ Code + Text

[ 65., 75.]]

T4 RAM Disk

Paused

↑ ↓ ↻ ⌨ ⚙ 📄 🗑 ⋮

▶

```
w = torch.randn(3,2, requires_grad=True)
b = torch.randn(2, requires_grad=True)

# Defining Model
def model(x):
    # return (x @ w) + b
    return torch.matmul(x,w)+b

#MSE Loss function
def mse(p1, p2):
    diff = p1-p2
    return torch.sum(diff**2)/ diff.numel()

# Training step
for i in range(1000):
    preds = model(inputs)
    loss = mse(preds, targets)
    loss.backward()
    if i%100 == 99:
        print(loss.item())

    with torch.no_grad():
        w -= w.grad * 0.00001
        b -= b.grad * 0.00001
        w.grad.zero_()
        b.grad.zero_()
```

AI SCIENCES

0s completed at 3:02 PM

Udemy

Pytorch.ipynb - Colab

Loss Functions.pdf

File

C:/Users/mg/Downloads/Loss%20Functions.pdf.pdf

Search

Star

Share

Download

Paused

Loss Functions.pdf

1 / 3 | - 70% + |

Download

Print

# Loss Functions

Loss Function	Use Case	Pros	Cons
Mean Squared Error (MSE)	Regression	Simple to compute and differentiate.	Sensitive to outliers, can lead to slow convergence.
Mean Absolute Error (MAE)	Regression	Robust to outliers.	Less smooth gradients compared to MSE, can be slower to converge.
Huber Loss	Regression (robust to outliers)	Balances sensitivity and robustness to outliers.	Requires tuning of the hyperparameter $\delta$ .
Cross-Entropy Loss	Classification (binary and multiclass)	Effective for classification tasks, especially with softmax output.	Can be sensitive to class imbalance.
Binary Cross-Entropy	Binary Classification	Suitable for binary classification, handles probabilities well.	Can suffer from vanishing gradients for extreme predictions.
Categorical Cross-Entropy	Multiclass Classification	Standard for multiclass classification with one-hot encoded labels.	Assumes mutually exclusive classes, not suitable for multi-label classification.
Sparse Categorical	Multiclass Classification with	Efficient for large number of classes, avoids one-hot	Similar issues as categorical cross-entropy

AI SCIENCES

Udemy

Pytorch.ipynb - Colab

Loss Functions.pdf

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=BZYFUI5XqzSc

+

Code

+

Text

All changes saved

T4

RAM

Disk

Paused

[ 65., 75.]]

```
w = torch.randn(3,2, requires_grad=True)
b = torch.randn(2, requires_grad=True)

# Defining Model
def model(x):
    # return (x @ w) + b
    return torch.matmul(x,w)+b

#MSE Loss function
def mse(p1, p2):
    diff = p1-p2
    return torch.sum(diff**2)/ diff.numel()

# Training step
for i in range(1000):
    preds = model(inputs)
    loss = mse(preds, targets)
    loss.backward()
    if i%100 == 99:
        print(loss.item())

    with torch.no_grad():
        w -= w.grad * 0.00001
        b -= b.grad * 0.00001
        w.grad.zero_()
        b.grad.zero_()
```

ΔSCIENCES

0s completed at 3:02 PM

UdenX

Pytorch.ipynb - Colab

Loss Functions.pdf

File

C:/Users/mg/Downloads/Loss%20Functions.pdf.pdf

Search

Star

Share

Download

Paused

Loss Functions.pdf

1 / 3 | - 70% + |

Download

Print

# Loss Functions

Loss Function	Use Case	Pros	Cons
Mean Squared Error (MSE)	Regression	Simple to compute and differentiate.	Sensitive to outliers, can lead to slow convergence.
Mean Absolute Error (MAE)	Regression	Robust to outliers.	Less smooth gradients compared to MSE, can be slower to converge.
Huber Loss	Regression (robust to outliers)	Balances sensitivity and robustness to outliers.	Requires tuning of the hyperparameter $\delta$ .
Cross-Entropy Loss	Classification (binary and multiclass)	Effective for classification tasks, especially with softmax output.	Can be sensitive to class imbalance.
Binary Cross-Entropy	Binary Classification	Suitable for binary classification, handles probabilities well.	Can suffer from vanishing gradients for extreme predictions.
Categorical Cross-Entropy	Multiclass Classification	Standard for multiclass classification with one-hot encoded labels.	Assumes mutually exclusive classes, not suitable for multi-label classification.
Sparse Categorical	Multiclass Classification with	Efficient for large number of classes, avoids one-hot	Similar issues as categorical cross-entropy

AI SCIENCES

Udemy



Pytorch.ipynb - Colab

Loss Functions.pdf

File

C:/Users/mg/Downloads/Loss%20Functions.pdf.pdf

Search

Star

Share

Download

Paused

Loss Functions.pdf

1 / 3 | - 70% +

Download

Print

# Loss Functions

Loss Function	Use Case	Pros	Cons
Mean Squared Error (MSE)	Regression	Simple to compute and differentiate.	Sensitive to outliers, can lead to slow convergence.
Mean Absolute Error (MAE)	Regression	Robust to outliers.	Less smooth gradients compared to MSE, can be slower to converge.
Huber Loss	Regression (robust to outliers)	Balances sensitivity and robustness to outliers.	Requires tuning of the hyperparameter $\delta$ .
Cross-Entropy Loss	Classification (binary and multiclass)	Effective for classification tasks, especially with softmax output.	Can be sensitive to class imbalance.
Binary Cross-Entropy	Binary Classification	Suitable for binary classification, handles probabilities well.	Can suffer from vanishing gradients for extreme predictions.
Categorical Cross-Entropy	Multiclass Classification	Standard for multiclass classification with one-hot encoded labels.	Assumes mutually exclusive classes, not suitable for multi-label classification.
Sparse Categorical	Multiclass Classification with	Efficient for large number of classes, avoids one-hot	Similar issues as categorical cross-entropy

AI SCIENCES

Udemy

Pytorch.ipynb - Colab

Loss Functions.pdf

File

C:/Users/mg/Downloads/Loss%20Functions.pdf.pdf

Search

Star

Share

Download

Paused

Loss Functions.pdf

1 / 3 | - 70% + |

Download

Print

# Loss Functions

Loss Function	Use Case	Pros	Cons
Mean Squared Error (MSE)	Regression	Simple to compute and differentiate.	Sensitive to outliers, can lead to slow convergence.
Mean Absolute Error (MAE)	Regression	Robust to outliers.	Less smooth gradients compared to MSE, can be slower to converge.
Huber Loss	Regression (robust to outliers)	Balances sensitivity and robustness to outliers.	Requires tuning of the hyperparameter $\delta$ .
Cross-Entropy Loss	Classification (binary and multiclass)	Effective for classification tasks, especially with softmax output.	Can be sensitive to class imbalance.
Binary Cross-Entropy	Binary Classification	Suitable for binary classification, handles probabilities well.	Can suffer from vanishing gradients for extreme predictions.
Categorical Cross-Entropy	Multiclass Classification	Standard for multiclass classification with one-hot encoded labels.	Assumes mutually exclusive classes, not suitable for multi-label classification.
Sparse Categorical	Multiclass Classification with	Efficient for large number of classes, avoids one-hot	Similar issues as categorical cross-entropy

AI SCIENCES

Udemy

Pytorch.ipynb - Colab

Loss Functions.pdf

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=BZYFUI5XqzSc

+

Code

+

Text

All changes saved

T4

RAM

Disk

Paused

[ 65., 75.]]

```
w = torch.randn(3,2, requires_grad=True)
b = torch.randn(2, requires_grad=True)

# Defining Model
def model(x):
    # return (x @ w) + b
    return torch.matmul(x,w)+b

#MSE Loss function
def mse(p1, p2):
    diff = p1-p2
    return torch.sum(diff**2)/ diff.numel()

# Training step
for i in range(1000):
    preds = model(inputs)
    loss = mse(preds, targets)
    loss.backward()
    if i%100 == 99:
        print(loss.item())

with torch.no_grad():
    w -= w.grad * 0.00001
    b -= b.grad * 0.00001
    w.grad.zero_()
    b.grad.zero_()
```

0s

completed at 3:02 PM

AI SCIENCES

Udacity

Pytorch.ipynb - Colab

Loss Functions.pdf

File

C:/Users/mg/Downloads/Loss%20Functions.pdf.pdf

Search

Star

Share

Download

Paused

Loss Functions.pdf

1 / 3 | - 70% +

Download

Print

More

# Loss Functions

Loss Function	Use Case	Pros	Cons
Mean Squared Error (MSE)	Regression	Simple to compute and differentiate.	Sensitive to outliers, can lead to slow convergence.
Mean Absolute Error (MAE)	Regression	Robust to outliers.	Less smooth gradients compared to MSE, can be slower to converge.
Huber Loss	Regression (robust to outliers)	Balances sensitivity and robustness to outliers.	Requires tuning of the hyperparameter $\delta$ .
Cross-Entropy Loss	Classification (binary and multiclass)	Effective for classification tasks, especially with softmax output.	Can be sensitive to class imbalance.
Binary Cross-Entropy	Binary Classification	Suitable for binary classification, handles probabilities well.	Can suffer from vanishing gradients for extreme predictions.
Categorical Cross-Entropy	Multiclass Classification	Standard for multiclass classification with one-hot encoded labels.	Assumes mutually exclusive classes, not suitable for multi-label classification.
Sparse Categorical	Multiclass Classification with	Efficient for large number of classes, avoids one-hot	Similar issues as categorical cross-entropy

AI SCIENCES

Udemy



Pytorch.ipynb - Colab

Loss Functions.pdf

File

C:/Users/mg/Downloads/Loss%20Functions.pdf.pdf

Search

Star

Share

Download

Paused

Loss Functions.pdf

1 / 3 | - 70% +

Download

Print

# Loss Functions

Loss Function	Use Case	Pros	Cons
Mean Squared Error (MSE)	Regression	Simple to compute and differentiate.	Sensitive to outliers, can lead to slow convergence.
Mean Absolute Error (MAE)	Regression	Robust to outliers.	Less smooth gradients compared to MSE, can be slower to converge.
Huber Loss	Regression (robust to outliers)	Balances sensitivity and robustness to outliers.	Requires tuning of the hyperparameter $\delta$ .
Cross-Entropy Loss	Classification (binary and multiclass)	Effective for classification tasks, especially with softmax output.	Can be sensitive to class imbalance.
Binary Cross-Entropy	Binary Classification	Suitable for binary classification, handles probabilities well.	Can suffer from vanishing gradients for extreme predictions.
Categorical Cross-Entropy	Multiclass Classification	Standard for multiclass classification with one-hot encoded labels.	Assumes mutually exclusive classes, not suitable for multi-label classification.
Sparse Categorical	Multiclass Classification with	Efficient for large number of classes, avoids one-hot	Similar issues as categorical cross-entropy

AI SCIENCES

Udemy

Pytorch.ipynb - Colab

Loss Functions.pdf

File

C:/Users/mg/Downloads/Loss%20Functions.pdf.pdf

Search

Star

Share

Download

Paused

Loss Functions.pdf

1 / 3 | - 70% +

Download

Print

More

Loss Function	Use Case	Pros	Cons
Mean Squared Error (MSE)	Regression	Simple to compute and differentiate.	Sensitive to outliers, can lead to slow convergence.
Mean Absolute Error (MAE)	Regression	Robust to outliers.	Less smooth gradients compared to MSE, can be slower to converge.
Huber Loss	Regression (robust to outliers)	Balances sensitivity and robustness to outliers.	Requires tuning of the hyperparameter $\delta$ .
Cross-Entropy Loss	Classification (binary and multiclass)	Effective for classification tasks, especially with softmax output.	Can be sensitive to class imbalance.
Binary Cross-Entropy	Binary Classification	Suitable for binary classification, handles probabilities well.	Can suffer from vanishing gradients for extreme predictions.
Categorical Cross-Entropy	Multiclass Classification	Standard for multiclass classification with one-hot encoded labels.	Assumes mutually exclusive classes, not suitable for multi-label classification.
Sparse Categorical Cross-Entropy	Multiclass Classification with integer labels	Efficient for large number of classes, avoids one-hot encoding.	Similar issues as categorical cross-entropy with class imbalance.
Hinge Loss	Support Vector Machines (SVMs)	Good for maximum margin classifiers,	Not differentiable at the margin, less commonly used in deep learning.

AI SCIENCES

Udemy

Google

Cross-EntropyLoss

I

X



All

Images

Videos

Shopping

News

Books

Web

More

Tools

Showing results for **Cross-Entropy Loss**

Search instead for **Cross-EntropyLoss**

اردو میں

In English


Cross-entropy, also known as logarithmic loss or log loss, is a popular loss function used in machine learning to measure the performance of a classification model. Namely, it measures the difference between the discovered probability distribution of a classification model and the predicted values.



DataCamp

<https://www.datacamp.com> > ... > Machine Learning

Cross-Entropy Loss Function in Machine Learning - DataCamp



$$\begin{aligned} L &= - \sum_{i=1}^2 t_i \log(p_i) \\ &= -[t_1 \log(p_1) + t_2 \log(p_2)] \\ &= -[t \log(p) + (1-t) \log(1-p)] \end{aligned}$$

truth value taking a value 0 or 1 and  $p_i$  is the Softmax probability. If we have two classes 1 and 0 we can have  $t_1 = 1$  and  $t_2 = 0$  then  $p_1 + p_2 = 1 \implies p_1 = 1 - p_2$ . For the convenience we let  $t_1 = t, t_2 = 1 - t, p_1 = p$  and  $p_2 = 1 - p$ .

About featured snippets Feedback



People also ask

Cross-entropy :

Udemy



Pytorch.ipynb - Colab

Loss Functions.pdf

cross-entropy loss formula - Go

File

C:/Users/mg/Downloads/Loss%20Functions.pdf.pdf

70%

3 / 3

Loss Functions.pdf

3 / 3

70%

+

-

Download

Print

More


• Mean Absolute Error (MAE): Measures the average absolute differences between predicted and actual values. Less sensitive to outliers compared to MSE.

• PyTorch Syntax: `torch.nn.L1Loss()`

• Huber Loss: A combination of MSE and MAE, providing robustness to outliers and smooth gradients.

• PyTorch Syntax: `torch.nn.SmoothL1Loss()`

• Cross-Entropy Loss: Measures the difference between two probability distributions, commonly used for classification tasks.



[WWW.AISCIENCES.IO](http://WWW.AISCIENCES.IO)

• PyTorch Syntax: `torch.nn.CrossEntropyLoss()`

• Binary Cross-Entropy: A special case of cross-entropy for binary classification problems.

• PyTorch Syntax: `torch.nn.BCELoss()`


• Categorical Cross-Entropy: Used for multiclass classification where each output class is one-hot encoded.

• PyTorch Syntax: `torch.nn.CrossEntropyLoss()`

• Sparse Categorical Cross-Entropy: Similar to categorical cross-entropy but uses integer labels for classes.

• PyTorch Syntax: `torch.nn.CrossEntropyLoss()` (with integer labels)

• Hinge Loss: Used primarily for training Support Vector Machines (SVMs)



Udemy



Pytorch.ipynb - Colab

Loss Functions.pdf

cross-entropy loss formula - Go

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=B6Ymnu54qzSb

+ Code + Text All changes saved

T4 RAM Disk

Paused

```
[ 65.,  /5.]]

w = torch.randn(3,2, requires_grad=True)
b = torch.randn(2, requires_grad=True)

# Defining Model
def model(x):
    # return (x @ w) + b
    return torch.matmul(x,w)+b

#MSE Loss function
def mse(p1, p2):
    diff = p1-p2
    return torch.sum(diff**2)/ diff.numel()

# Training step
for i in range(1000):
    preds = model(inputs)
    loss = mse(preds, targets)
    loss.backward()
    if i%100 == 99:
        print(loss.item())

    with torch.no_grad():
        w -= w.grad * 0.00001
        b -= b.grad * 0.00001
        w.grad.zero_()
        b.grad.zero_()

preds = model(inputs)
```

0s completed at 3:02 PM

AI SCIENCES

Udacity

Pytorch.ipynb - Colab

Loss Functions.pdf

cross-entropy loss formula - Go

File

C:/Users/mg/Downloads/Loss%20Functions.pdf.pdf

Search

Star

Share

Download

Paused

Loss Functions.pdf

2 / 3

136%

Fullscreen

Refresh

	(imbalanced data)	hard examples.	$y \setminus \gamma$ may, requires tuning.
--	-------------------	----------------	--

- Mean Squared Error (MSE): Measures the average squared differences between predicted and actual values. Widely used in regression tasks.
  - PyTorch Syntax: `torch.nn.MSELoss()`
- Mean Absolute Error (MAE): Measures the average absolute differences between predicted and actual values. Less sensitive to outliers compared to MSE.
  - PyTorch Syntax: `torch.nn.L1Loss()`
- Huber Loss: A combination of MSE and MAE, providing robustness to outliers and smooth gradients.

AI SCIENCES

Udemy

```
w = torch.randn(3,2, requires_grad=True)
b = torch.randn(2, requires_grad=True)
```

```
def model(x):  
    # return (x @ w) + b  
    return torch.matmul(x,w)+b
```

```
#Replacing MSE with built-in func
```

```
for i in range(1000):
    preds = model(inputs)
    loss = mse(preds, targets)
    loss.backward()
    if i%100 == 99:
        print(loss.item())
```

```
w -= w.grad * 0.00001
b -= b.grad * 0.00001
w.grad.zero_()
b.grad.zero_()
```

```
preds = model(inputs)
loss = mse(preds, targets)
print(loss)
```

## Using Pytorch

The torch package contains data structures for multi-dimensional tensors and defines mathematical operations over these tensors. Additionally, it provides many utilities for efficient serialization of Tensors and arbitrary types, and other useful utilities.

It has a CUDA counterpart, that enables you to run your tensor computations on an NVIDIA GPU with compute capability  $\geq 3.0$ .

[illegible]



Pytorch.ipynb - Colab

Loss Functions.pdf

cross-entropy loss formula - Go

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=B6Ymnu54qzSb

☆

Paused

+ Code

+ Text

All changes saved

T4

RAM

Disk

Basic Neural Network

Using Pytorch

import torch

import numpy as np

import torch.nn.functional as F

# Input (temp, rainfall, humidity)

inputs = np.array([[73, 67, 43],

[91, 88, 64],

[87, 134, 58],

[102, 43, 37],

[69, 96, 70],

[85, 100, 60],

[95, 80, 55],

[105, 120, 75],

[78, 90, 50],

[82, 70, 45]], dtype='float32')

# Targets (apples, oranges)

targets = np.array([[56, 70],

[81, 101],

[119, 133],

[22, 37],

torch

The torch package contains data structures for multi-dimensional tensors and defines mathematical operations over these tensors. Additionally, it provides many utilities for efficient serialization of Tensors and arbitrary types, and other useful utilities.

It has a CUDA counterpart, that enables you to run your tensor computations on an NVIDIA GPU with compute capability >= 3.0.

↑

↓

ASCIENCES

0s

completed at 3:02 PM

Pytorch.ipynb - Colab

Loss Functions.pdf

cross-entropy loss formula - Go

File

C:/Users/mg/Downloads/Loss%20Functions.pdf.pdf

Search

Star

Share

Download

Paused

Loss Functions.pdf

2 / 3 | 102%

Fullscreen

Refresh

		(e.g., embeddings).	value prediction.
Focal Loss	Object detection and classification (imbalanced data)	Addresses class imbalance by focusing on hard examples.	Introduces an additional hyperparameter $\gamma$ , requires tuning.

- Mean Squared Error (MSE): Measures the average squared differences between predicted and actual values. Widely used in regression tasks.
  - PyTorch Syntax: `torch.nn.MSELoss()`
- Mean Absolute Error (MAE): Measures the average absolute differences between predicted and actual values. Less sensitive to outliers compared to MSE.
  - PyTorch Syntax: `torch.nn.L1Loss()`
- Huber Loss: A combination of MSE and MAE, providing robustness to outliers and smooth gradients.
  - PyTorch Syntax: `torch.nn.SmoothL1Loss()`
- Cross-Entropy Loss: Measures the difference between two probability

AI SCIENCES

Udemy



```
with torch.no_grad():
    w -= w.grad * 0.00001
    b -= b.grad * 0.00001
    w.grad.zero_()
    b.grad.zero_()
```

0s

↑↓



W

#

de

re

#

ms

#

fo

pm

10

10

it

DI

W

W

b

W

b

PM







```
def model(x):
```

```
# return (x @ w) + b
```

```
mse = torch.nn.MSELoss()
```

```
for i in range(1000):
```

```
preds = model(inputs)
```

```
loss = mse(preds, targets)
```

```
loss.backward()
```

```
if i%100 == 99:
```

```
print(loss.item())
```

```
w -= w.grad * 0.00001
```

```
b -= b.grad * 0.00001
```

w.grad.zero ()

```
b.grad.zero_()
```

```
preds = model(inputs)
```

Pytorch.ipynb - Colab

Loss Functions.pdf

cross-entropy loss formula - Go

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=BZYFUI5XqzSc

T4 RAM Disk

Paused

+ Code + Text

0s

```
for i in range(1000):
    preds = model(inputs)
    loss = mse(preds, targets)
    loss.backward()
    if i%100 == 99:
        print(loss.item())

    with torch.no_grad():
        w -= w.grad * 0.00001
        b -= b.grad * 0.00001
        w.grad.zero_()
        b.grad.zero_()

preds = model(inputs)
loss = mse(preds, targets)
print(loss)
```

↕

24.0968017578125  
183.68312072753906  
148.4185028076172  
127.21119689941406  
113.1971664428711  
103.01649475097656  
95.0095443725586  
88.34284973144531  
82.5859146118164  
77.50598907470703  
tensor(77.4581, grad\_fn=<MseLossBackward0>)

[ ]

AISCIENCES

completed at 3:10 PM

Pytorch.ipynb - Colab

Loss Functions.pdf

cross-entropy loss formula - Go

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=BZYFUI5XqzSc

☆

Paused

+ Code + Text

T4 RAM Disk

0s

[11]

[103., 119.],

[ 98., 110.],

[ 88., 95.],

[115., 140.],

[ 76., 85.],

[ 65., 75.]]

0s

w = torch.randn(3,2, requires\_grad=True)

b = torch.randn(2, requires\_grad=True)

# Defining Model

def model(x):

# return (x @ w) + b

return torch.matmul(x,w)+b

#Replacing MSE with built-in function

mse = torch.nn.MSELoss()

# Training step

for i in range(1000):

preds = model(inputs)

loss = mse(preds, targets)

loss.backward()

if i%100 == 99:

print(loss.item())

with torch.no\_grad():

w -= w.grad \* 0.00001

b -= b.grad \* 0.00001

w.grad.zero\_()

completed at 3:10 PM

AI SCIENCES

Udemy



Pytorch.ipynb - Colab

Loss Functions.pdf

cross-entropy loss formula - Go

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=BZYFUI5XqzSc

T4 RAM Disk

Paused

+ Code + Text

```
# Defining Model
def model(x):
    # return (x @ w) + b
    return torch.matmul(x,w)+b

#Replacing MSE with built-in function
# mse = torch.nn.MSELoss()
mae = torch.nn.L1Loss()

# Training step
for i in range(1000):
    preds = model(inputs)
    loss = mse(preds, targets)
    loss.backward()
    if i%100 == 99:
        print(loss.item())

    with torch.no_grad():
        w -= w.grad * 0.00001
        b -= b.grad * 0.00001
        w.grad.zero_()
        b.grad.zero_()

preds = model(inputs)
loss = mse(preds, targets)
print(loss)
```

247.0968017578125

183.68312072753906

0s

completed at 3:10 PM

AI SCIENCES

Udacity