

IMPORTANT ELEMENTS OF NEURAL NETWORK


- Loss Function



- Optimizer




- Activation Function





IMPORTANT ELEMENTS OF NEURAL NETWORK



- Loss Function ✓
- Optimizer
- Activation Function





~~SIMPLE~~ NEURAL NETWORK





~~SIMPLE~~ NEURAL NETWORK





~~SIMPLE~~ NEURAL NETWORK



~~SIMPLE~~ NEURAL NETWORK



~~SIMPLE~~ NEURAL NETWORK



Pytorch.ipynb - Colab

Activation Function.docx.pdf

relu vs leaky relu - Google Search

ReLU activation function vs. Leaky ReLU

File

C:/Users/mg/Downloads/Activation%20Function.docx.pdf.pdf

53%

2 / 4

Softmax	Output layers for multiclass classification	Converts logits to probabilities; good for classification	computationally intensive for many classes
Softplus	Hidden layers, smooth activation	Smooth approximation of ReLU; always positive	Computationally more complex than ReLU
Softsign	Hidden layers, alternative to Tanh	Bounded outputs; reduces vanishing gradient problem	Slower convergence compared to ReLU and Tanh

Descriptions

1. **Sigmoid:** Maps the input to a value between 0 and 1. Often used in the output layer of binary classification problems.

- PyTorch Syntax: `torch.nn.Sigmoid()`

2. **Tanh:** Maps the input to a value between -1 and 1. Often used in hidden layers and can be seen as a scaled version of the sigmoid function.

- PyTorch Syntax: `torch.nn.Tanh()`

3. **ReLU (Rectified Linear Unit):** Applies the rectifier function, which outputs the input directly if it is positive; otherwise, it outputs zero. Commonly used in hidden layers of neural networks.

- PyTorch Syntax: `torch.nn.ReLU()`

4. **Leaky ReLU:** A variation of ReLU that allows a small, non-zero gradient when the input is negative, which helps to avoid dying ReLUs.

- PyTorch Syntax: `torch.nn.LeakyReLU()`

5. **ELU (Exponential Linear Unit):** An activation function that tends to converge faster and more reliably than ReLU by smoothing the transition from negative to positive inputs.

- PyTorch Syntax: `torch.nn.ELU()`

6. **SELU (Scaled Exponential Linear Unit):** A self-normalizing activation

AI SCIENCES

Udemy

Pytorch.ipynb - Colab

Activation Function.docx.pdf

relu vs leaky relu - Google Search

ReLU activation function vs. Leaky ReLU

File

C:/Users/mg/Downloads/Activation%20Function.docx.pdf.pdf

Search

Star

Share

Download

Paused

Activation Function.docx.pdf

1 / 4 | - 53% +

Fullscreen

Refresh

Download

Print

More

Activation Functions

Activation Function	Use Case	Pros	Cons
Sigmoid	Binary classification, output layers	Outputs between 0 and 1; good for probabilities	Vanishing gradient problem
Tanh	Hidden layers, regression tasks	Outputs between -1 and 1; zero-centered	Vanishing gradient problem
ReLU (Rectified Linear Unit)	Hidden layers, convolutional neural networks (CNNs)	Computationally efficient; reduces vanishing gradient	Can cause dying ReLUs (neurons stop learning)
Leaky ReLU	Hidden layers, avoiding dying ReLUs	Prevents dying ReLUs by allowing small gradients	Introduces a small slope for negative inputs
ELU (Exponential Linear Unit)	Hidden layers, deep neural networks	Smooths the transition, faster and reliable convergence	More computationally expensive
SELU (Scaled Exponential Linear Unit)	Self-normalizing networks, hidden layers	Self-normalizes outputs; robust training	Requires specific initialization and dropout scaling
Softmax	Output layers for multiclass classification	Converts logits to probabilities; good for classification	Can be computationally intensive for many classes

AI SCIENCES

Udemy

Pytorch.ipynb - Colab

Activation Function.docx.pdf

relu vs leaky relu - Google Search

ReLU activation function vs. Leaky ReLU

File

C:/Users/mg/Downloads/Activation%20Function.docx.pdf.pdf

Search

Star

Share

Download

Paused

Activation Function.docx.pdf

1 / 4 | - 53% +

Fullscreen

Refresh

Download

Print

More

Activation Functions

Activation Function	Use Case	Pros	Cons
Sigmoid	Binary classification, output layers	Outputs between 0 and 1; good for probabilities	Vanishing gradient problem
Tanh	Hidden layers, regression tasks	Outputs between -1 and 1; zero-centered	Vanishing gradient problem
ReLU (Rectified Linear Unit)	Hidden layers, convolutional neural networks (CNNs)	Computationally efficient; reduces vanishing gradient	Can cause dying ReLUs (neurons stop learning)
Leaky ReLU	Hidden layers, avoiding dying ReLUs	Prevents dying ReLUs by allowing small gradients	Introduces a small slope for negative inputs
ELU (Exponential Linear Unit)	Hidden layers, deep neural networks	Smooths the transition, faster and reliable convergence	More computationally expensive
SELU (Scaled Exponential Linear Unit)	Self-normalizing networks, hidden layers	Self-normalizes outputs; robust training	Requires specific initialization and dropout scaling
Softmax	Output layers for multiclass classification	Converts logits to probabilities; good for classification	Can be computationally intensive for many classes

AI SCIENCES

Udemy

Pytorch.ipynb - Colab

Activation Function.docx.pdf

relu vs leaky relu - Google Search

ReLU activation function vs. Leaky ReLU

File

C:/Users/mg/Downloads/Activation%20Function.docx.pdf.pdf

Search

Star

Share

Download

Paused

Activation Function.docx.pdf

1 / 4

53%

Fullscreen

Refresh

Download

Print

More

Function			
Sigmoid	Binary classification, output layers	Outputs between 0 and 1; good for probabilities	Vanishing gradient problem
Tanh	Hidden layers, regression tasks	Outputs between -1 and 1; zero-centered	Vanishing gradient problem
ReLU (Rectified Linear Unit)	Hidden layers, convolutional neural networks (CNNs)	Computationally efficient; reduces vanishing gradient	Can cause dying ReLUs (neurons stop learning)
Leaky ReLU	Hidden layers, avoiding dying ReLUs	Prevents dying ReLUs by allowing small gradients	Introduces a small slope for negative inputs
ELU (Exponential Linear Unit)	Hidden layers, deep neural networks	Smooths the transition, faster and reliable convergence	More computationally expensive
SELU (Scaled Exponential Linear Unit)	Self-normalizing networks, hidden layers	Self-normalizes outputs; robust training	Requires specific initialization and dropout scaling
Softmax	Output layers for multiclass classification	Converts logits to probabilities; good for classification	Can be computationally intensive for many classes
Softplus	Hidden layers, smooth activation	Smooth approximation of ReLU; always positive	Computationally more complex than ReLU

AI SCIENCES

Udemy

Pytorch.ipynb - Colab

Activation Function.docx.pdf

relu vs leaky relu - Google Search

ReLU activation function vs. Leaky ReLU

File

C:/Users/mg/Downloads/Activation%20Function.docx.pdf.pdf

Search

Star

Share

Download

Paused

Activation Function.docx.pdf

1 / 4

53%

Fullscreen

Refresh

Download

Print

More

Function			
Sigmoid	Binary classification, output layers	Outputs between 0 and 1; good for probabilities	Vanishing gradient problem
Tanh	Hidden layers, regression tasks	Outputs between -1 and 1; zero-centered	Vanishing gradient problem
ReLU (Rectified Linear Unit)	Hidden layers, convolutional neural networks (CNNs)	Computationally efficient; reduces vanishing gradient	Can cause dying ReLUs (neurons stop learning)
Leaky ReLU	Hidden layers, avoiding dying ReLUs	Prevents dying ReLUs by allowing small gradients	Introduces a small slope for negative inputs
ELU (Exponential Linear Unit)	Hidden layers, deep neural networks	Smooths the transition, faster and reliable convergence	More computationally expensive
SELU (Scaled Exponential Linear Unit)	Self-normalizing networks, hidden layers	Self-normalizes outputs; robust training	Requires specific initialization and dropout scaling
Softmax	Output layers for multiclass classification	Converts logits to probabilities; good for classification	Can be computationally intensive for many classes
Softplus	Hidden layers, smooth activation	Smooth approximation of ReLU; always positive	Computationally more complex than ReLU

AI SCIENCES

Udemy

Google

relu vs leaky relu

All Images Videos News Shopping Books Web More

Tools

Saved



Sigmoid



Softmax



Prelu activation



Neural networks



Function




Linear



Swish activation




Parametric



ResearchGate

ReLU activation function vs. LeakyReLU ...




Medium

Activation functions: ReLU vs. Leaky ...

ResearchGate

ResearchGate



ReLU activation function vs. LeakyReLU activation function. |...

Visit >

Images may be subject to copyright. Learn More

Share

Save

AI SCIENCES







Figure 3 - available via license: [Creative Commons Attribution-NonCommercial-NoDerivatives 4.0 International](#)
Content may be subject to copyright.

[Download](#) [View publication](#) [v](#)




ReLU activation function vs. LeakyReLU activation function.

[Source publication](#)


Figure 3 - available via license: [Creative Commons Attribution-NonCommercial-NoDerivatives 4.0 International](#)
Content may be subject to copyright.

Download

View publication



ReLU activation function



LeakyReLU activation function

ReLU activation function vs. LeakyReLU activation function.

[Source publication](#)

Pytorch.ipynb - Colab

Activation Function.docx.pdf

relu vs leaky relu - Google Search

ReLU activation function vs. Leaky ReLU

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=BZYFUI5XqzSc

☆

Paused

+ Code

+ Text

All changes saved

0s

[76., 85.],
[65., 75.]]

#Replacing Model with Built-in Model

class SimpleNN(nn.Module):

#Replacing MSE with built-in function

mse = torch.nn.MSELoss()

mae = torch.nn.L1Loss()

Training step

for i in range(1000):

preds = model(inputs)

loss = mae(preds, targets)

loss.backward()

if i%100 == 99:

print(loss.item())

with torch.no_grad():

w -= w.grad * 0.00001

b -= b.grad * 0.00001

w.grad.zero_()

b.grad.zero_()

preds = model(inputs)

loss = mae(preds, targets)

print(loss)

126.0269546508789

AI SCIENCES

0s completed at 3:11 PM

Pytorch.ipynb - Colab

Activation Function.docx.pdf

relu vs leaky relu - Google Search

ReLU activation function vs. Leaky ReLU

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=BZYFUI5XqzSc

Paused

T4 RAM Disk

0s

[76., 85.],
[65., 75.]]

#Replacing Model with Built-in Model

cal

#Replacing MSE with built-in function

mse = torch.nn.MSELoss()

mae = torch.nn.L1Loss()

Training step

for i in range(1000):

preds = model(inputs)

loss = mae(preds, targets)

loss.backward()

if i%100 == 99:

print(loss.item())

with torch.no_grad():

w -= w.grad * 0.00001

b -= b.grad * 0.00001

w.grad.zero_()

b.grad.zero_()

preds = model(inputs)

loss = mae(preds, targets)

print(loss)

126.0269546508789

0s completed at 3:11 PM

AISCIENCES

Udemy

Pytorch.ipynb - Colab

Activation Function.docx.pdf

relu vs leaky relu - Google Search

ReLU activation function vs. Leaky ReLU

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=BZYFUI5XqzSc

✓ T4 RAM Disk

Paused

+ Code

+ Text

All changes saved

✓ 0s

[11]

[103., 119.],
[98., 110.],
[88., 95.],
[115., 140.],
[76., 85.],
[65., 75.]]

▶

w = torch.randn(3,2, requires_grad=True)
b = torch.randn(2, requires_grad=True)

Defining Model
def model(x):
 # return (x @ w) + b
 return torch.matmul(x,w)+b

#Replacing Model with built-in function
class SimpleNN(nn.Module):
 def __init__(self):
 super().__init__()

#Replacing MSE with built-in function
mse = torch.nn.MSELoss()
mae = torch.nn.L1Loss()

Training step
for i in range(1000):

I

ASCIENCES

0s completed at 3:11 PM

Udacity

[↑]

```
[ 21.,  57.],  
[103., 119.],  
[ 98., 110.],  
[ 88.,  95.],  
[115., 140.],  
[ 76.,  85.],  
[ 65.,  75.]])
```

```
w = torch.randn(3,2, requires_grad=True)
b = torch.randn(2, requires_grad=True)
```

```
# Defining Model
```

```
def model(x):
```

```
# return (x @ w) + b
```

```
return torch.matmul(x,w)+b
```

#Replacing Model with built-in function

```
class SimpleNN(nn.Module):
```

```
def __init__(self):
```

```
super().__init__()
```

```
self.linear1 = nn.Linear()
```

#Replacing MSE with built-in function

```
# mse = torch.nn.MSELoss()
```

```
mae = torch.nn.L1Loss()
```



```
# Training step
```

```
for i in range(1000):
```



Δ SCIENCES

~~SIMPLE~~ NEURAL NETWORK



Pytorch.ipynb - Colab

Activation Function.docx.pdf

relu vs leaky relu - Google Search

ReLU activation function vs. Leaky ReLU

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=BZYFUI5XqzSc

0s

✓

T4

RAM

Disk

Paused

+ Code

+ Text

All changes saved

✓

0s

▶

↕

[103., 119.],

[98., 110.],

[88., 95.],

[115., 140.],

[76., 85.],

[65., 75.]]

▶

w = torch.randn(3,2, requires_grad=True)

b = torch.randn(2, requires_grad=True)

Defining Model

def model(x):

return (x @ w) + b

return torch.matmul(x,w)+b

#Replacing Model with built-in function

class SimpleNN(nn.Module):

def __init__(self):

super().__init__()

self.linear1 = nn.Linear(3, 4)

self.act1

self.linear2 = nn.Linear(4, 2)

#Replacing MSE with built-in function

mse = torch.nn.MSELoss()

mae = torch.nn.L1Loss()

↑

↓

🔗

💬

⚙️

📄

🗑️

⋮

ASCIENCES

0s

completed at 3:11 PM

Udacity

Pytorch.ipynb - Colab

Activation Function.docx.pdf

relu vs leaky relu - Google Search

ReLU activation function vs. Leaky ReLU

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=BZYFUI5XqzSc

✓ T4 RAM Disk

Paused

+ Code + Text All changes saved

✓ 0s

[22., 37.],

[103., 119.],

[98., 110.],

[88., 95.],

[115., 140.],

[76., 85.],

[65., 75.]]

▶

w = torch.randn(3,2, requires_grad=True)

b = torch.randn(2, requires_grad=True)

Defining Model

def model(x):

return (x @ w) + b

return torch.matmul(x,w)+b

#Replacing Model with built-in function

class SimpleNN(nn.Module):

def __init__(self):

super().__init__()

self.linear1 = nn.Linear(3, 4)

self.act1 = nn.ReLU

self.linear2 = nn.Linear(4, 2)

#Replacing MSE with built-in function

mse = torch.nn.MSELoss()

mae = torch.nn.L1Loss()



↑ ↓ ↻ ⌨ ⚙ 📄 🗑 ⋮

AI SCIENCES

0s completed at 3:11 PM

UdenX

~~SIMPLE~~ NEURAL NETWORK



Pytorch.ipynb - Colab

Activation Function.docx.pdf

relu vs leaky relu - Google Search

ReLU activation function vs. Leaky ReLU

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=BZYFUI5XqzSc

☆

Paused

+ Code + Text All changes saved

T4 RAM Disk

↑ ↓ ↻ ⌨ ⚙ 📄 🗑 ⋮

[65., 75.]]

▶

```
w = torch.randn(3,2, requires_grad=True)
b = torch.randn(2, requires_grad=True)

# Defining Model
def model(x):
    # return (x @ w) + b
    return torch.matmul(x,w)+b

#Replacing Model with built-in function
class SimpleNN(nn.Module):
    def __init__(self):
        super().__init__()
        self.linear1 = nn.Linear(3, 4)
        self.act1 = nn.ReLU()
        self.linear2 = nn.Linear(4, 2)

    def forward(self, x):
        x = self.linear1(x)
        x = self.act1(x)
        x = self.linear2(x)
        return x

#Replacing MSE with built-in function
# mse = torch.nn.MSELoss()
mae = torch.nn.L1Loss()
```

<>

ASCIENCES

0s completed at 3:11 PM

Udemy


```
w = torch.randn(3,2, requires_grad=True)
b = torch.randn(2, requires_grad=True)
```

```
def model(x):
```

```
# return  $(x @ w) + b$ 
return torch.matmul(x,w)+b
```

```
class SimpleNN(nn.Module):
```

```
super().__init__(
```

```
super().__init__()
```

```
... self.linear1 = nn.Linear(3, 4)
```

```
... self.act1 = nn.ReLU()
```

```
self.linear2 = nn.Linear(4, 2)
```

```
x = self.linear1(x)
```

```
x = self.linear1(x)
```

```
x = self.act1(x)
```

```
x = self.linear2(x)
```

```
return x
```

✓ 0s completed at 3:11 PM



```
w = torch.randn(3,2, requires_grad=True)
b = torch.randn(2, requires_grad=True)
```

```
def model(x):
```

```
# return (x @ w) + b
return torch.matmul(x,w)+b
```

```
class SimpleNN(nn.Module):
```

```
def __init__(self):
    super().__init__()
    self.linear1 = nn.Linear(3, 4)
    self.act1 = nn.ReLU()
    self.linear2 = nn.Linear(4, 2)
```

```
def forward(self, x):
    x = self.linear1(x)
    x = self.act1(x)
    x = self.linear2(x)
    return x
```

✓ 0s completed at 3:11 PM

Pytorch.ipynb - Colab

Activation Function.docx.pdf

relu vs leaky relu - Google Search

ReLU activation function vs. Leaky ReLU

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=BZYFUI5XqzSc

Paused

Reconnect

Settings

+ Code + Text Saving...

[65., 75.]]

▶

#Replacing Model with built-in function

class SimpleNN(nn.Module):

def __init__(self):

super().__init__()

self.linear1 = nn.Linear(3, 4)

self.act1 = nn.ReLU()

self.linear2 = nn.Linear(4, 2)

def forward(self, x):

x = self.linear1(x)

x = self.act1(x)

x = self.linear2(x)

return x

#Replacing MSE with built-in function

mse = torch.nn.MSELoss()

mae = torch.nn.L1Loss()

Training step

for i in range(1000):

preds = model(inputs)

0s completed at 3:11 PM

AI SCIENCES

Udacity

Pytorch.ipynb - Colab

Activation Function.docx.pdf

relu vs leaky relu - Google Search

ReLU activation function vs. Leaky ReLU

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=BZYFUI5XqzSc

RAM
Disk

Paused

+ Code + Text

```
## Training step
# for i in range(1000):
#     preds = model(inputs)
#     loss = mae(preds, targets)
#     loss.backward()
#     if i%100 == 99:
#         print(loss.item())

#         with torch.no_grad():
#             w -= w.grad * 0.00001
#             b -= b.grad * 0.00001
#             w.grad.zero_()
#             b.grad.zero_()

#     preds = model(inputs)
#     loss = mae(preds, targets)
#     print(loss)
```

↔

126.0269546508789
116.77412414550781
107.52127838134766
98.26844787597656
89.01560974121094
80.09945678710938
72.49869537353516
66.56382751464844
61.56400680541992
57.83696365356445
tensor(57.8033, grad_fn=<MeanBackward0>)

7s completed at 3:43 PM

AI SCIENCES

Udemy

Pytorch.ipynb - Colab

Activation Function.docx.pdf

relu vs leaky relu - Google Search

ReLU activation function vs. Leaky ReLU

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=BZYFUI5XqzSc

Paused

+ Code + Text

class SimpleNN(nn.Module):
 def __init__(self):
 super().__init__()
 self.linear1 = nn.Linear(3, 4)
 self.act1 = nn.ReLU()
 self.linear2 = nn.Linear(4, 2)

 def forward(self, x):
 x = self.linear1(x)
 x = self.act1(x)
 x = self.linear2(x)
 return x

model = SimpleNN()

#Replacing MSE with built-in function
mse = torch.nn.MSELoss()
mae = torch.nn.L1Loss()

Training step
for i in range(1000):
preds = model(inputs)
loss = mae(preds, targets)
loss.backward()
if i%100 == 99:
print(loss.item())

with torch.no_grad():

RAM

Disk

↑ ↓ ↻ ⌨ ⚙ 📄 🗑 ⋮

7s completed at 3:43 PM

ASCIENCES

Udemy

Pytorch.ipynb - Colab

Activation Function.docx.pdf

relu vs leaky relu - Google Search

ReLU activation function vs. Leaky ReLU

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=BZYFUI5XqzSc

☆

Paused

RAM

Disk

↑

↓

↻

💬

⚙️

📄

🗑️

⋮

+ Code

+ Text

All changes saved

▶

#Replacing Model with built-in function

class SimpleNN(nn.Module):

def __init__(self):

super().__init__()

self.linear1 = nn.Linear(3, 4)

self.act1 = nn.ReLU()

self.linear2 = nn.Linear(4, 2)

def forward(self, x):

x = self.linear1(x)

x = self.act1(x)

x = self.linear2(x)

return x

model = SimpleNN()

for name, params in model.named_parameters():

print("Name:", name)

print("Param",)

#Replacing MSE with built-in function

mse = torch.nn.MSELoss()

mae = torch.nn.L1Loss()

Training step

for i in range(1000):

preds = model(inputs)

loss = mae(preds, targets)

7s

completed at 3:43 PM

Udacity

Pytorch.ipynb - Colab

Activation Function.docx.pdf

relu vs leaky relu - Google Search

ReLU activation function vs. Leaky ReLU

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=B6Ymnu54qzSb

☆

Paused

+ Code + Text All changes saved

✓ RAM Disk

Basic Neural Network with Built-in functions

Using Pytorch

7s

▶

```
import torch
import numpy as np
import torch.nn.functional as F

# Input (temp, rainfall, humidity)
inputs = np.array([[73, 67, 43],
                  [91, 88, 64],
                  [87, 134, 58],
                  [102, 43, 37],
                  [69, 96, 70],
                  [85, 100, 60],
                  [95, 80, 55],
                  [105, 120, 75],
                  [78, 90, 50],
                  [82, 70, 45]], dtype='float32')

# Targets (apples, oranges)
targets = np.array([[56, 70],
                   [81, 101],
                   [119, 133],
                   [22, 37],
                   [103, 119],
```

0s completed at 3:44 PM

AI SCIENCES

Udacity

Pytorch.ipynb - Colab

Activation Function.docx.pdf

relu vs leaky relu - Google Search

ReLU activation function vs. Leaky ReLU

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=B6Ymnu54qzSb

☆

Paused

+ Code + Text

✓ RAM Disk

Basic Neural Network with Built-in functions

Using Pytorch

import torch
import numpy as np
import torch.nn as nn

Input (temp, rainfall, humidity)
inputs = np.array([[73, 67, 43],
[91, 88, 64],
[87, 134, 58],
[102, 43, 37],
[69, 96, 70],
[85, 100, 60],
[95, 80, 55],
[105, 120, 75],
[78, 90, 50],
[82, 70, 45]], dtype='float32')

Targets (apples, oranges)
targets = np.array([[56, 70],
[81, 101],
[119, 133],
[22, 37],
[103, 119],

0s completed at 3:44 PM

AI SCIENCES


```
+ Code + Text All changes saved
# 1e-05 == 1e-05.
# print(loss.item())

# with torch.no_grad():
#     w -= w.grad * 0.00001
#     b -= b.grad * 0.00001
#     w.grad.zero_()
#     b.grad.zero_()

# preds = model(inputs)
# loss = mae(preds, targets)
# print(loss)
```

```
↔ Name: linear1.weight
Param Parameter containing:
tensor([[ -0.1478, -0.5037, -0.5325],
        [ 0.3457,  0.3467, -0.5086],
        [-0.1229, -0.5641, -0.5510],
        [-0.3043,  0.1261, -0.0497]], requires_grad=True)
Name: linear1.bias
Param Parameter containing:
tensor([0.4177, 0.4991, 0.3199, 0.0991], requires_grad=True)
Name: linear2.weight
Param Parameter containing:
tensor([[ 0.3285,  0.4916, -0.0986,  0.4325],
        [ 0.4481,  0.4910, -0.3399, -0.2481]], requires_grad=True)
Name: linear2.bias
Param Parameter containing:
tensor([0.4516, 0.2715], requires_grad=True)
```

[]

Pytorch.ipynb - Colab

Activation Function.docx.pdf

relu vs leaky relu - Google Search

ReLU activation function vs. Leaky ReLU

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=BZYFUI5XqzSc

RAM

Disk

Paused

+ Code + Text All changes saved

0s

```
# 1e100 == inf.
# print(loss.item())

# with torch.no_grad():
#     w -= w.grad * 0.00001
#     b -= b.grad * 0.00001
#     w.grad.zero_()
#     b.grad.zero_()

# preds = model(inputs)
# loss = mae(preds, targets)
# print(loss)
```

↔

Name: linear1.weight

Param Parameter containing:

tensor([[-0.1478, -0.5037, -0.5325],
 [0.3457, 0.3467, -0.5086],
 [-0.1229, -0.5641, -0.5510],
 [-0.3043, 0.1261, -0.0497]], requires_grad=True)

Name: linear1.bias

Param Parameter containing:

tensor([0.4177, 0.4991, 0.3199, 0.0991], requires_grad=True)

Name: linear2.weight

Param Parameter containing:

tensor([[0.3285, 0.4916, -0.0986, 0.4325],
 [0.4481, 0.4910, -0.3399, -0.2481]], requires_grad=True)

Name: linear2.bias

Param Parameter containing:

tensor([0.4516, 0.2715], requires_grad=True)

[]

0s completed at 3:45 PM

AI SCIENCES

Udacity

Pytorch.ipynb - Colab

Activation Function.docx.pdf

relu vs leaky relu - Google Search

ReLU activation function vs. Leaky ReLU

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=BZYFUI5XqzSc

RAM

Disk

Paused

+ Code

+ Text

All changes saved

0s

b.grad.zero_()

preds = model(inputs)

loss = mae(preds, targets)

print(loss)

→

Name: linear1.weight

Param Parameter containing:

tensor([[-0.1478, -0.5037, -0.5325],

[0.3457, 0.3467, -0.5086],

[-0.1229, -0.5641, -0.5510],

[-0.3043, 0.1261, -0.0497]], requires_grad=True)

Name: linear1.bias

Param Parameter containing:

tensor([0.4177, 0.4991, 0.3199, 0.0991], requires_grad=True)

Name: linear2.weight

Param Parameter containing:

tensor([[0.3285, 0.4916, -0.0986, 0.4325],

[0.4481, 0.4910, -0.3399, -0.2481]], requires_grad=True)

Name: linear2.bias

Param Parameter containing:

tensor([0.4516, 0.2715], requires_grad=True)

[]

[]

[]

completed at 3:45 PM

AI SCIENCES

Udemy

✓

↑↓

Pytorch.ipynb - Colab

Activation Function.docx.pdf

relu vs leaky relu - Google Search

ReLU activation function vs. Leaky ReLU

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=BZYFUI5XqzSc

RAM

Disk

Paused

+ Code + Text All changes saved

0s

```
# b.grad.zero_()

# preds = model(inputs)
# loss = mae(preds, targets)
# print(loss)
```

↔

Name: linear1.weight

Param Parameter containing:

tensor([[-0.1478, -0.5037, -0.5325],

[0.3457, 0.3467, -0.5086],

[-0.1229, -0.5641, -0.5510],

[-0.3043, 0.1261, -0.0497]], requires_grad=True)

Name: linear1.bias

Param Parameter containing:

tensor([0.4177, 0.4991, 0.3199, 0.0991], requires_grad=True)

Name: linear2.weight

Param Parameter containing:

tensor([[0.3285, 0.4916, -0.0986, 0.4325],

[0.4481, 0.4910, -0.3399, -0.2481]], requires_grad=True)

Name: linear2.bias

Param Parameter containing:

tensor([0.4516, 0.2715], requires_grad=True)

[]

[]

[]

ASCIENCES

0s completed at 3:45 PM

Udemy

Pytorch.ipynb - Colab

Activation Function.docx.pdf

relu vs leaky relu - Google Search

ReLU activation function vs. Leaky ReLU

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=BZYFUI5XqzSc

☆

Paused

RAM

Disk

+ Code

+ Text

All changes saved

0s

[3]

[98., 110.],

[88., 95.],

[115., 140.],

[76., 85.],

[65., 75.]])

#Replacing Model with built-in function

class SimpleNN(nn.Module):

def __init__(self):

super().__init__()

self.linear1 = nn.Linear(3, 4)

self.act1 = nn.ReLU()

self.linear2 = nn.Linear(4, 2)

def forward(self, x):

x = self.linear1(x)

x = self.act1(x)

x = self.linear2(x)

return x

model = SimpleNN()

for name, params in model.named_parameters():

print("Name:", name)

print("Param", params)

#Replacing MSE with built-in function

mse = torch.nn.MSELoss()

mae = torch.nn.L1Loss()

0s

completed at 3:45 PM

AISCIENCES

Udacity

Pytorch.ipynb - Colab

Activation Function.docx.pdf

relu vs leaky relu - Google Search

ReLU activation function vs. Leaky ReLU

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=BZYFUI5XqzSc

RAM
Disk

Paused

+ Code

+ Text

All changes saved

0s

[3]

[115., 140.],
[76., 85.],
[65., 75.]]

0s

#Replacing Model with built-in function

class SimpleNN(nn.Module):

def __init__(self):

super().__init__()

self.linear1 = nn.Linear(3, 4)

self.act1 = nn.ReLU()

self.linear2 = nn.Linear(4, 2)

def forward(self, x):

x = self.linear1(x)

x = self.act1(x)

x = self.linear2(x)

return x

model = SimpleNN()

for name, params in model.named_parameters():

print("Name:", name)

print("Param", params)

#Replacing MSE with built-in function

mse = torch.nn.MSELoss()

mae = torch.nn.L1Loss()

ASCIENCES

WWW.ASCIENCES.ACADEMY/COURSE-PYTORCH

0s completed at 3:45 PM

Udemy

Pytorch.ipynb - Colab

Activation Function.docx.pdf

relu vs leaky relu - Google Search

ReLU activation function vs. Leaky ReLU

colab.research.google.com/drive/1imQs4frEn2sYP0VW7SjxgJwC5ZgNZdpv#scrollTo=BZYFUI5XqzSc

Paused

RAM
Disk

+ Code + Text All changes saved

0s [3] [115., 140.],
[76., 85.],
[65., 75.]]

#Replacing Model with built-in function

class SimpleNN(nn.Module):

def __init__(self):

super().__init__()

self.linear1 = nn.Linear(3, 4)

self.act1 = nn.ReLU()

self.linear2 = nn.Linear(4, 2)

def forward(self, x):

x = self.linear1(x)

x = self.act1(x)

x = self.linear2(x)

return x

model = SimpleNN()

for name, params in model.named_parameters():

print("Name:", name)

print("Param", params)

#Replacing MSE with built-in function

mse = torch.nn.MSELoss()

mae = torch.nn.L1Loss()

WWW.ASCIENCES.ACADEMY/COURSE-PYTORCH

completed at 3:45 PM

ASCIENCES

Udacity